

# Assignment1, COMP 576

Yikun Li (yl212)

September 29, 2021

## Task 1.a

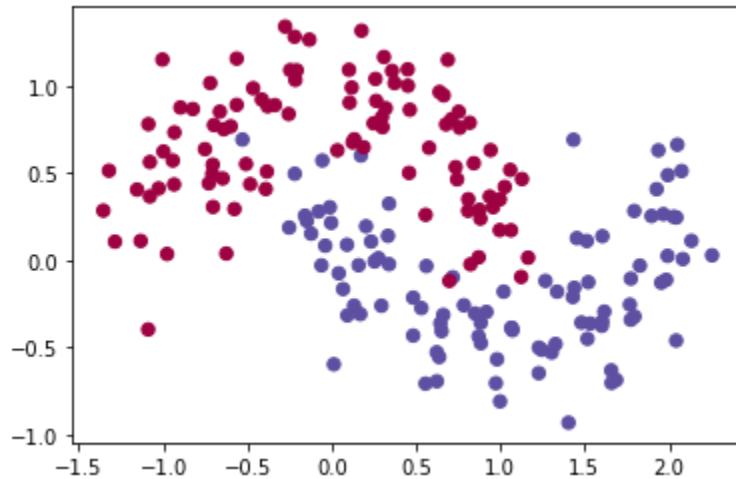


Figure 1\_1: generate\_data

### 1.b.1

```
def actFun(self, z, type):
    """
    actFun computes the activation functions
    :param z: net input
    :param type: Tanh, Sigmoid, or ReLU
    :return: activations
    """

    # YOU IMPLEMENT YOUR actFun HERE
    if type == "tanh":
        return np.tanh(z)
    elif type == "sigmoid":
        return 1 / (1 + np.exp(-z))
    elif type == "relu":
        return np.maximum(0, z)
    else:
        return None
```

## 1.b.2

Derive the derivatives of Tanh, Sigmoid and ReLU:

$$\frac{d}{dx} \tanh(x) = 1 - \tanh^2(x)$$

$$\frac{d}{dx} \sigma(x) = \sigma(x)(1 - \sigma(x))$$

$$\frac{d}{x} \text{relu}(x) = x > 0 ? 1 : 0$$

## 1.b.3

```
def diff_actFun(self, z, type):
    """
    diff_actFun compute the derivatives of the activation functions wrt the net input
    :param z: net input
    :param type: Tanh, Sigmoid, or ReLU
    :return: the derivatives of the activation functions wrt the net input
    """

    # YOU IMPLEMENT YOUR diff_actFun HERE
    if type == 'tanh':
        return 1 - np.square(np.tanh(z))
    elif type == 'sigmoid':
        tmp = 1 / (1 + np.exp(-z))
        return tmp * (1 - tmp)
    elif type == 'relu':
        return np.where(z > 0, 1, 0)
    else:
        return None
```

### 1.c.1

```
def feedforward(self, X, actFun):
    """
    feedforward builds a 3-layer neural network and computes the two probabilities,
    one for class 0 and one for class 1
    :param X: input data
    :param actFun: activation function
    :return:
    """

    # YOU IMPLEMENT YOUR feedforward HERE

    self.z1 = np.dot(X, self.W1) + self.b1
    self.a1 = actFun(self.z1)
    self.z2 = np.dot(self.a1, self.W2) + self.b2
    exp_scores = np.exp(self.z2)
    self.probs = exp_scores / np.sum(exp_scores, axis=1, keepdims=True)
    return None
```

### 1.c.2

```
def calculate_loss(self, X, y):
    """
    calculate_loss compute the loss for prediction
    :param X: input data
    :param y: given labels
    :return: the loss for prediction
    """

    num_examples = len(X)
    self.feedforward(X, lambda x: self.actFun(x, type=self.actFun_type))
    # Calculating the loss

    # YOU IMPLEMENT YOUR CALCULATION OF THE LOSS HERE

    probs = np.exp(self.z2) / np.sum(np.exp(self.z2), axis=1, keepdims=True)
    data_loss_single = -np.log(probs[range(num_examples), y])
    data_loss = np.sum(data_loss_single)

    # Add regularization term to loss (optional)
    data_loss += self.reg_lambda / 2 * (np.sum(np.square(self.W1)) + np.sum(np.square(self.W2)))
    return (1. / num_examples) * data_loss
```

## 1.d.1

$$\Delta_3 = \frac{\partial L}{\partial z_2} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z_2}$$

$$\frac{\partial L}{\partial W_2} = a_1^T \Delta_3$$

$$\frac{\partial L}{\partial b_2} = \sum \Delta_3$$

$$\Delta_2 = diff * \Delta_3 W_2^T$$

$$\frac{\partial L}{\partial W_1} = X^T \Delta_2$$

$$\frac{\partial L}{\partial b_2} = \sum \Delta_2$$

## 1.d.2

```
def backprop(self, X, y):
    """
    backprop run backpropagation to compute the gradients used to update the parameters in the backward step
    :param X: input data
    :param y: given labels
    :return: dL/dW1, dL/b1, dL/dW2, dL/db2
    """

    # IMPLEMENT YOUR BACKPROP HERE
    num_examples = len(X)
    delta3 = self.probs
    delta3[range(num_examples), y] -= 1
    dW2 = np.dot(self.a1.T, delta3)
    db2 = np.sum(delta3, axis=0, keepdims=True)
    diff = self.diff_actFun(self.z1, type=self.actFun_type)
    delta2 = np.dot(delta3, self.W2.T) * diff
    dW1 = np.dot(X.T, delta2)
    db1 = np.sum(delta2, axis=0, keepdims=False)
    return dW1, dW2, db1, db2
```

### 1.e.1: Change Type

Tanh:

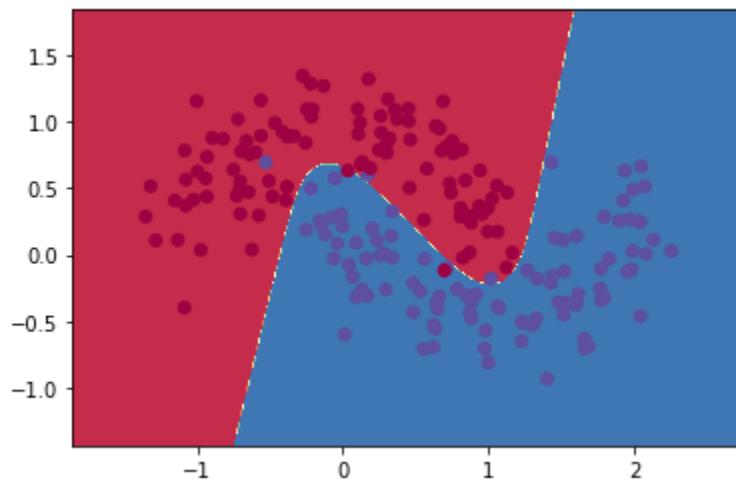


Figure 1\_2: Tanh

Sigmoid:

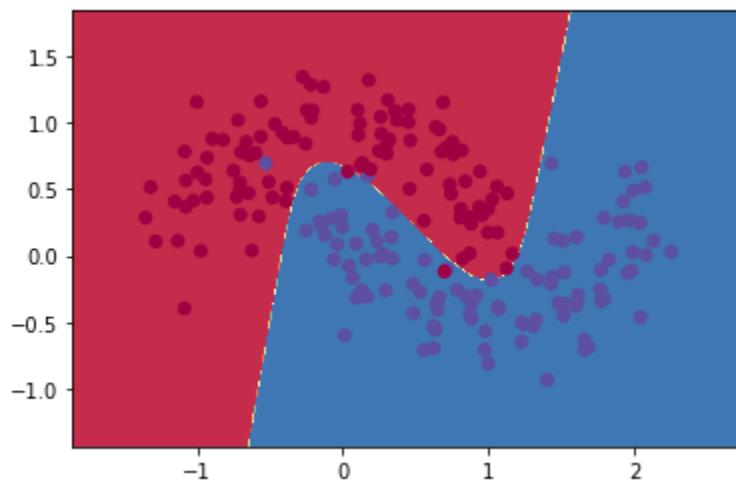


Figure 1\_3: Sigmoid

ReLU:

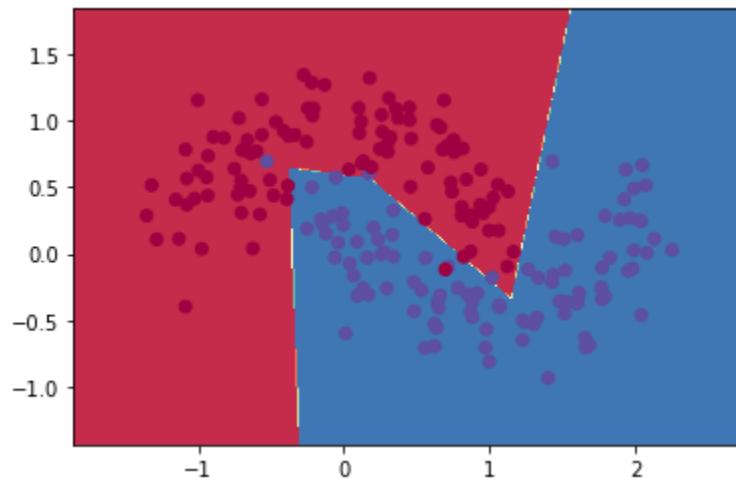


Figure 1\_4: ReLu

Observation and explanation: ReLu performs better than Tanh and Sigmoid, the latter two performing very similarly. The reason for this situation is that the ReLu can make some of the output to be zero which can lead to Scale-invariant.

### 1.e.2: Change nn\_hidden\_dim

Type = Tanh, nn\_hidden\_dim = 1

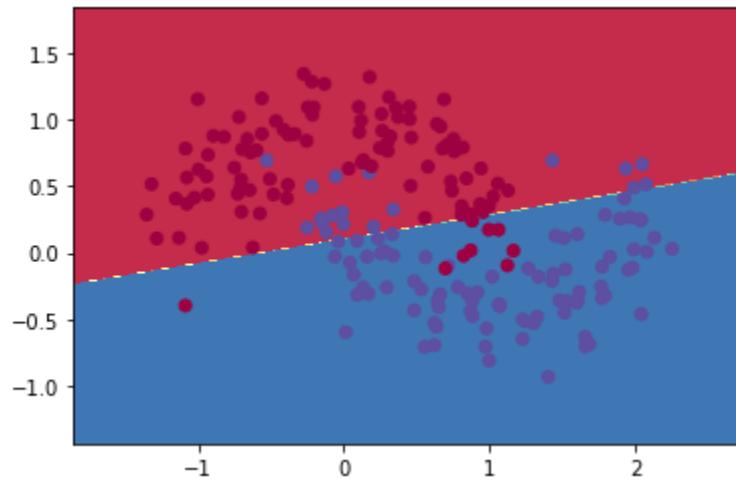


Figure 1\_5: nn\_hidden\_dim = 1

Type = Tanh, nn\_hidden\_dim = 3

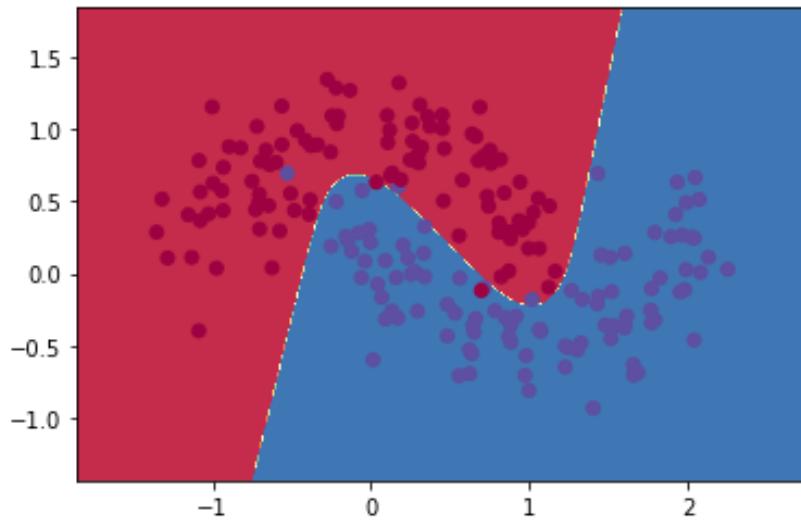


Figure 1\_6: nn\_hidden\_dim = 3

Type = Tanh, nn\_hidden\_dim = 5

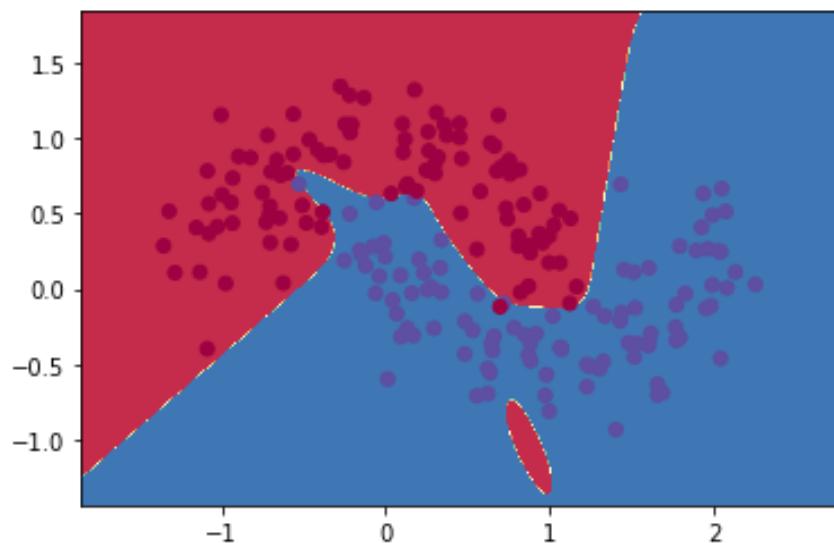


Figure 1\_7: nn\_hidden\_dim = 5

Type = Tanh, nn\_hidden\_dim = 7

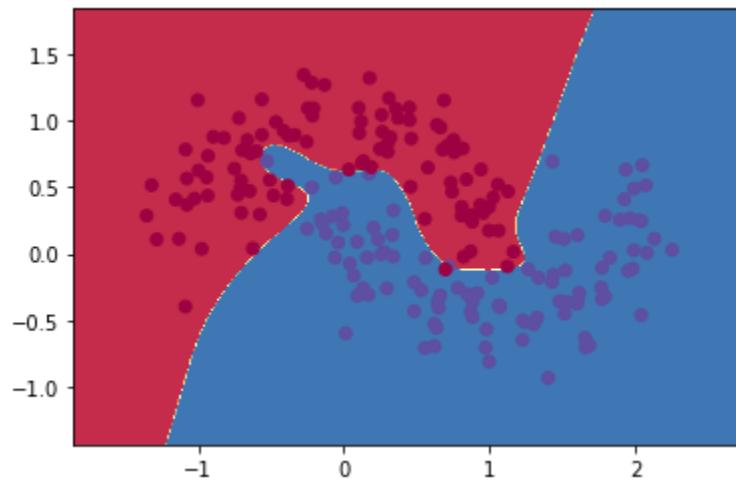


Figure 1\_8: nn\_hidden\_dim = 7

Type = Tanh, nn\_hidden\_dim = 9

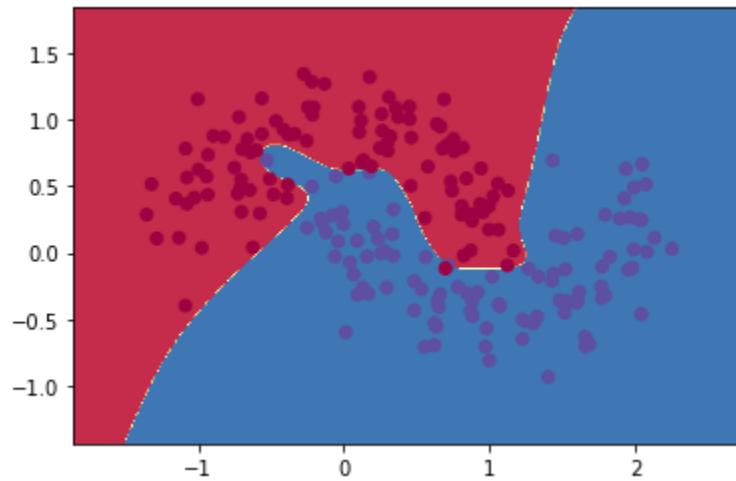


Figure 1\_9: nn\_hidden\_dim = 9

Observation: When the nn\_hidden\_dim = 3, the performance is the best. Because if the value is too small, the model will unfit the data; if the value is too big, the model will overfit the data.

## 1.f : DeepNeuralNetwork implementation

In my implementation, call the DeepNeuralNetwork this:

```
model = DeepNeuralNetwork(nn_dims=[2, 3, 2], actFun_type='tanh')
```

Which means the model is constructed by 3 layers with 2, 3, 2 units on each layer from input to output, and the active function is tanh.

Here are some detailed implementation on key functions:

Function feedforward:

```
def feedforward(self, X, actFun):
    """
    feedforward builds a 3-layer neural network and computes the two probabilities,
    one for class 0 and one for class 1
    :param X: input data
    :param actFun: activation function
    :return:
    """

# YOU IMPLEMENT YOUR feedforward HERE

    self.z = []
    self.a = []
    for i in range(len(self.W)):
        if i == 0:
            self.z.append(np.dot(X, self.W[i]) + self.b[i])
        else:
            self.z.append(np.dot(self.a[i-1], self.W[i]) + self.b[i])
        if i != len(self.W) - 1:
            self.a.append(actFun(self.z[i]))
    exp_scores = np.exp(self.z[len(self.z)-1])
    self.probs = exp_scores / np.sum(exp_scores, axis=1, keepdims=True)
    return None
```

Function calculate\_loss:

```
def calculate_loss(self, X, y):
    """
    calculate_loss computes the loss for prediction
    :param X: input data
    :param y: given labels
    :return: the loss for prediction
    """
    num_examples = len(X)
    self.feedforward(X, lambda x: self.actFun(x, type=self.actFun_type))

    # YOU IMPLEMENT YOUR CALCULATION OF THE LOSS HERE

    probs = np.exp(self.z[len(self.z)-1]) / \
            np.sum(np.exp(self.z[len(self.z)-1]), axis=1, keepdims=True)
    data_loss_single = -np.log(probs[range(num_examples), y])
    data_loss = np.sum(data_loss_single)

    # Add regularization term to loss (optional)
    W_sum = 0
    for i in len(self.W):
        W_sum += np.sum(np.square(self.W[i]))
    data_loss += self.reg_lambda / 2 * W_sum
    return (1. / num_examples) * data_loss
```

Function backprop:

```
def backprop(self, X, y):
    """
    backprop implements backpropagation to compute the gradients used to update the parameters in the backward step
    :param X: input data
    :param y: given labels
    :return: dL/dW1, dL/b1, dL/dW2, dL/db2, ... dL/dn, dL/bn in two lists
    """
    # IMPLEMENT YOUR BACKPROP HERE
    num_examples = len(X)
    delta = self.probs
    delta[range(num_examples), y] -= 1

    dW = []
    db = []
    for i in range(len(self.z)):
        index = len(self.z) - i - 1
        if index != 0:
            dW.insert(0, np.dot(self.a[index - 1].T, delta))
            db.insert(0, np.sum(delta, axis=0, keepdims=True))
            delta = np.dot(delta, self.W[index].T) * \
                self.diff_actFun(self.z[index-1], type=self.actFun_type)
        else:
            dW.insert(0, np.dot(X.T, delta))
            db.insert(0, np.sum(delta, axis=0, keepdims=False))

    return dW, db
```

Function fit\_model:

```

def fit_model(self, X, y, epsilon=0.01, num_passes=20000, print_loss=True):
    """
    fit_model uses backpropagation to train the network
    :param X: input data
    :param y: given labels
    :param num_passes: the number of times that the algorithm runs through the whole dataset
    :param print_loss: print the loss or not
    :return:
    """
    # Gradient descent.
    for i in range(0, num_passes):
        # Forward propagation
        self.feedforward(X, lambda x: self.actFun(x, type=self.actFun_type))
        # Backpropagation
        dW, db = self.backprop(X, y)

        # Add regularization terms (b1 and b2 don't have regularization terms)
        for i in range(len(dW)):
            # print(dW[i].shape)
            # print(self.W[i].shape)
            dW[i] += self.reg_lambda * self.W[i]

        # Gradient descent parameter update
        for i in range(len(self.W)):
            self.W[i] += -epsilon * dW[i]
            self.b[i] += -epsilon * db[i]

        # Optionally print the loss.
        # This is expensive because it uses the whole dataset, so we don't want to do it too often.
        if print_loss and i % 1000 == 0:
            print("Loss after iteration %i: %f" % (i, self.calculate_loss(X, y)))

```

## 1.f: Change number of layers

In this section, let's keep the dataset to be make\_moons and the active function to be tanh. The change part is the number of layers and each time we add a three-unit layer in the network. Let's observe the difference on the boundary of images.

When we set the layer as:

```
model = DeepNeuralNetwork(nn_dims=[2, 3, 2], actFun_type='tanh')
```

The result is:

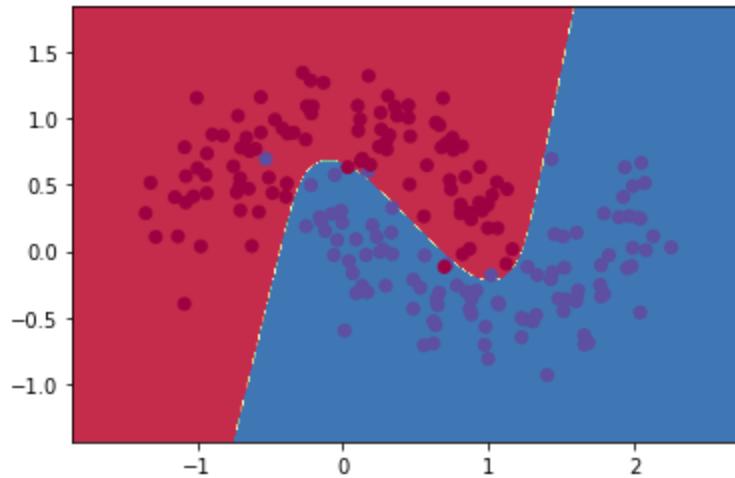


Figure 1\_10: nn\_hidden\_dim = 3

When we set the layer as:

```
model = DeepNeuralNetwork(nn_dims=[2, 3, 3, 2], actFun_type='tanh')
```

The result is:

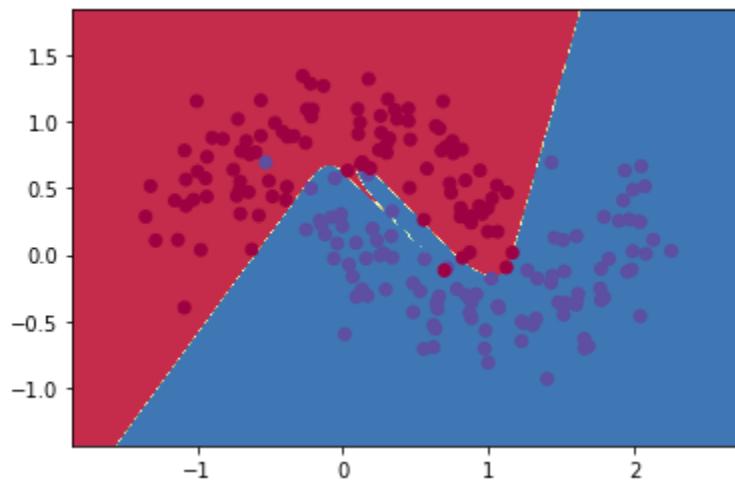


Figure 1\_11: nn\_hidden\_dim = 4

When we set the layer as:

```
model = DeepNeuralNetwork(nn_dims=[2, 3, 3, 3, 2], actFun_type='tanh')
```

The result is:

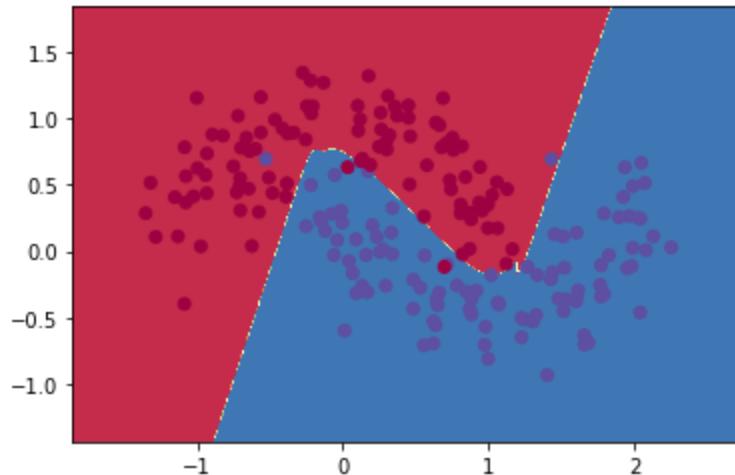


Figure 1\_12: nn\_hidden\_dim = 5

Observation: When the number of layers is 3 or 5, the performance is better than the number of layers to be 4. The layers of the network go deeper, the performance may be better, although this is not a must.

## 1.f: Change dataset

When the dataset is changed to be make\_circles:

```
X, y = datasets.make_circles(  
    n_samples=100, shuffle=True, noise=None, random_state=None, factor=0.8)
```

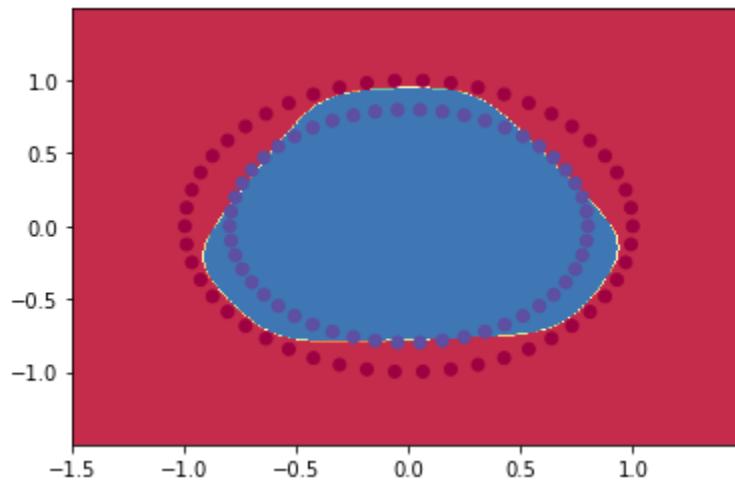


Figure 1\_13: dataset=make\_circles

This dataset is made up of data points in circles and are classified into two classes, with each of them making up a circle. As the boundary shows in the above image, our three-layer network can classify it well.

## Task 2.a.2

Reference from the tutorial GitHub repo of tensorflow:

<https://gist.github.com/saitodev/c4c7a8c83f5aa4a00e93084dd3f848c5>

```
def weight_variable(shape):
    """
    Initialize weights
    :param shape: shape of weights, e.g. [w, h ,cin, cout] where
    w: width of the filters
    h: height of the filters
    cin: the number of the channels of the filters
    cout: the number of filters
    :return: a tensor variable for weights with initial values
    """

    # IMPLEMENT YOUR WEIGHT_VARIABLE HERE
    initial = tf.truncated_normal(shape, stddev=0.1)
    return tf.Variable(initial)

def bias_variable(shape):
    """
    Initialize biases
    :param shape: shape of biases, e.g. [cout] where
    cout: the number of filters
    :return: a tensor variable for biases with initial values
    """

    # IMPLEMENT YOUR BIAS_VARIABLE HERE
    initial = tf.constant(0.1, shape=shape)
    return tf.Variable(initial)
```

```

def conv2d(x, W):
    """
    Perform 2-D convolution
    :param x: input tensor of size [N, W, H, Cin] where
    N: the number of images
    W: width of images
    H: height of images
    Cin: the number of channels of images
    :param W: weight tensor [w, h, Cin, Cout]
    w: width of the filters
    h: height of the filters
    Cin: the number of the channels of the filters = the number of channels of images
    Cout: the number of filters
    :return: a tensor of features extracted by the filters, a.k.a. the results after convolution
    """

    # IMPLEMENT YOUR CONV2D HERE
    return tf.nn.conv2d(x, W, strides=[1, 1, 1, 1], padding='SAME')


def max_pool_2x2(x):
    """
    Perform non-overlapping 2-D maxpooling on 2x2 regions in the input data
    :param x: input data
    :return: the results of maxpooling (max-marginalized + downsampling)
    """

    # IMPLEMENT YOUR MAX_POOL_2X2 HERE
    return tf.nn.max_pool(x, ksize=[1, 2, 2, 1],
                          strides=[1, 2, 2, 1], padding='SAME')

```

## 2.a.3

Reference from the tutorial GitHub repo of tensorflow:  
<https://gist.github.com/saitodev/c4c7a8c83f5aa4a00e93084dd3f848c5>

```
# placeholders for input data and input labeles
x = tf.placeholder(tf.float32, [None, 784], name='x')
y_ = tf.placeholder(tf.float32, [None, 10], name='y_')

# reshape the input image
x_image = tf.reshape(x, [-1, 28, 28, 1])

# first convolutional layer
W_conv1 = weight_variable([5, 5, 1, 32])
b_conv1 = bias_variable([32])
h_conv1 = tf.nn.relu(conv2d(x_image, W_conv1) + b_conv1)
h_pool1 = max_pool_2x2(h_conv1)
variable_summaries(W_conv1)
variable_summaries(b_conv1)
variable_summaries(h_conv1)
variable_summaries(h_pool1)

# second convolutional layer
W_conv2 = weight_variable([5, 5, 32, 64])
b_conv2 = bias_variable([64])
h_conv2 = tf.nn.relu(conv2d(h_pool1, W_conv2) + b_conv2)
h_pool2 = max_pool_2x2(h_conv2)
variable_summaries(W_conv2)
variable_summaries(b_conv2)
variable_summaries(h_conv2)
variable_summaries(h_pool2)

# densely connected layer
W_fc1 = weight_variable([7 * 7 * 64, 1024])
b_fc1 = bias_variable([1024])
h_pool2_flat = tf.reshape(h_pool2, [-1, 7*7*64])
h_fc1 = tf.nn.relu(tf.matmul(h_pool2_flat, W_fc1) + b_fc1)
variable_summaries(W_fc1)
variable_summaries(b_fc1)
variable_summaries(h_pool2_flat)
variable_summaries(h_fc1)
```

```

# dropout
keep_prob = tf.placeholder(tf.float32)
h_fc1_drop = tf.nn.dropout(h_fc1, keep_prob)
variable_summaries(keep_prob)
variable_summaries(h_fc1_drop)

# softmax
W_fc2 = weight_variable([1024, 10])
b_fc2 = bias_variable([10])
y_conv = tf.nn.softmax(tf.matmul(h_fc1_drop, W_fc2) + b_fc2, name='y_conv')
variable_summaries(W_fc2)
variable_summaries(b_fc2)
variable_summaries(y_conv)

```

## 2.a.4

Reference from the tutorial GitHub repo of tensorflow:

<https://gist.github.com/saitodev/c4c7a8c83f5aa4a00e93084dd3f848c5>

```

# setup training
cross_entropy = tf.reduce_mean(-tf.reduce_sum(y_ * tf.log(y_conv), reduction_indices=[1]))
train_step = tf.train.AdamOptimizer(1e-4).minimize(cross_entropy)
correct_prediction = tf.equal(tf.argmax(y_conv, 1), tf.argmax(y_, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32), name='accuracy')

```

## 2.a.5

```
# run the training
for i in range(max_step):
    batch = mnist.train.next_batch(50) # make the data batch, which is used in the training iteration.
    # the batch size is 50
    if i % 100 == 0:
        # output the training accuracy every 100 iterations
        train_accuracy = accuracy.eval(feed_dict={
            x: batch[0], y_: batch[1], keep_prob: 1.0})
        print("step %d, training accuracy %g" % (i, train_accuracy))

        # Update the events file which is used to monitor the training (in this case,
        # only the training loss is monitored)
        summary_str = sess.run(summary_op, feed_dict={x: batch[0], y_: batch[1], keep_prob: 0.5})
        summary_writer.add_summary(summary_str, i)
        summary_writer.flush()

    # save the checkpoints every 1100 iterations
    if i % 1100 == 0 or i == max_step:
        checkpoint_file = os.path.join(result_dir, 'checkpoint')
        saver.save(sess, checkpoint_file, global_step=i)

    train_step.run(feed_dict={x: batch[0], y_: batch[1], keep_prob: 0.5}) # run one train_step

# print test error
print("test accuracy %g" % accuracy.eval(feed_dict={
    x: mnist.test.images, y_: mnist.test.labels, keep_prob: 1.0}))

stop_time = time.time()
print('The training takes %f second to finish' % (stop_time - start_time))
```

Here is the final output results:

step 0, training accuracy 0.08  
step 100, training accuracy 0.9  
step 200, training accuracy 0.96  
step 300, training accuracy 0.96  
step 400, training accuracy 0.84  
step 500, training accuracy 0.92  
step 600, training accuracy 0.94  
step 700, training accuracy 0.9  
step 800, training accuracy 0.94  
step 900, training accuracy 0.98  
step 1000, training accuracy 0.98  
step 1100, training accuracy 0.96  
step 1200, training accuracy 0.96

step 1300, training accuracy 1  
step 1400, training accuracy 0.96  
step 1500, training accuracy 0.98  
step 1600, training accuracy 0.96  
step 1700, training accuracy 0.96  
step 1800, training accuracy 0.98  
step 1900, training accuracy 0.96  
step 2000, training accuracy 0.98  
step 2100, training accuracy 0.98  
step 2200, training accuracy 1  
step 2300, training accuracy 1  
step 2400, training accuracy 0.94  
step 2500, training accuracy 0.96  
step 2600, training accuracy 1  
step 2700, training accuracy 0.98  
step 2800, training accuracy 1  
step 2900, training accuracy 0.96  
step 3000, training accuracy 1  
step 3100, training accuracy 1  
step 3200, training accuracy 0.98  
step 3300, training accuracy 0.96  
step 3400, training accuracy 1  
step 3500, training accuracy 0.98  
step 3600, training accuracy 1  
step 3700, training accuracy 1  
step 3800, training accuracy 1  
step 3900, training accuracy 0.96  
step 4000, training accuracy 1  
step 4100, training accuracy 1  
step 4200, training accuracy 1  
step 4300, training accuracy 1  
step 4400, training accuracy 0.98  
step 4500, training accuracy 0.98  
step 4600, training accuracy 1  
step 4700, training accuracy 1  
step 4800, training accuracy 0.98  
step 4900, training accuracy 1  
step 5000, training accuracy 0.98  
step 5100, training accuracy 0.96  
step 5200, training accuracy 0.98

step 5300, training accuracy 0.98  
step 5400, training accuracy 1  
test accuracy 0.9864  
The training takes 812.483530 second to finish

## 2.a.6

There are three tags in the TensorBoard: scalars, graphs, and projector.

The figure of scalars:

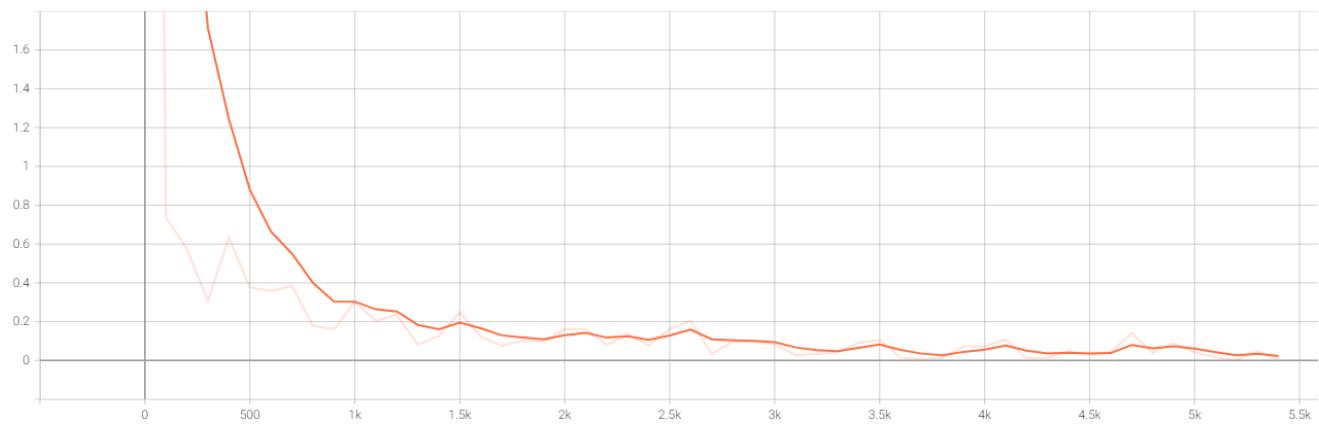


Figure 2\_1: scalars

The figure of graphs:

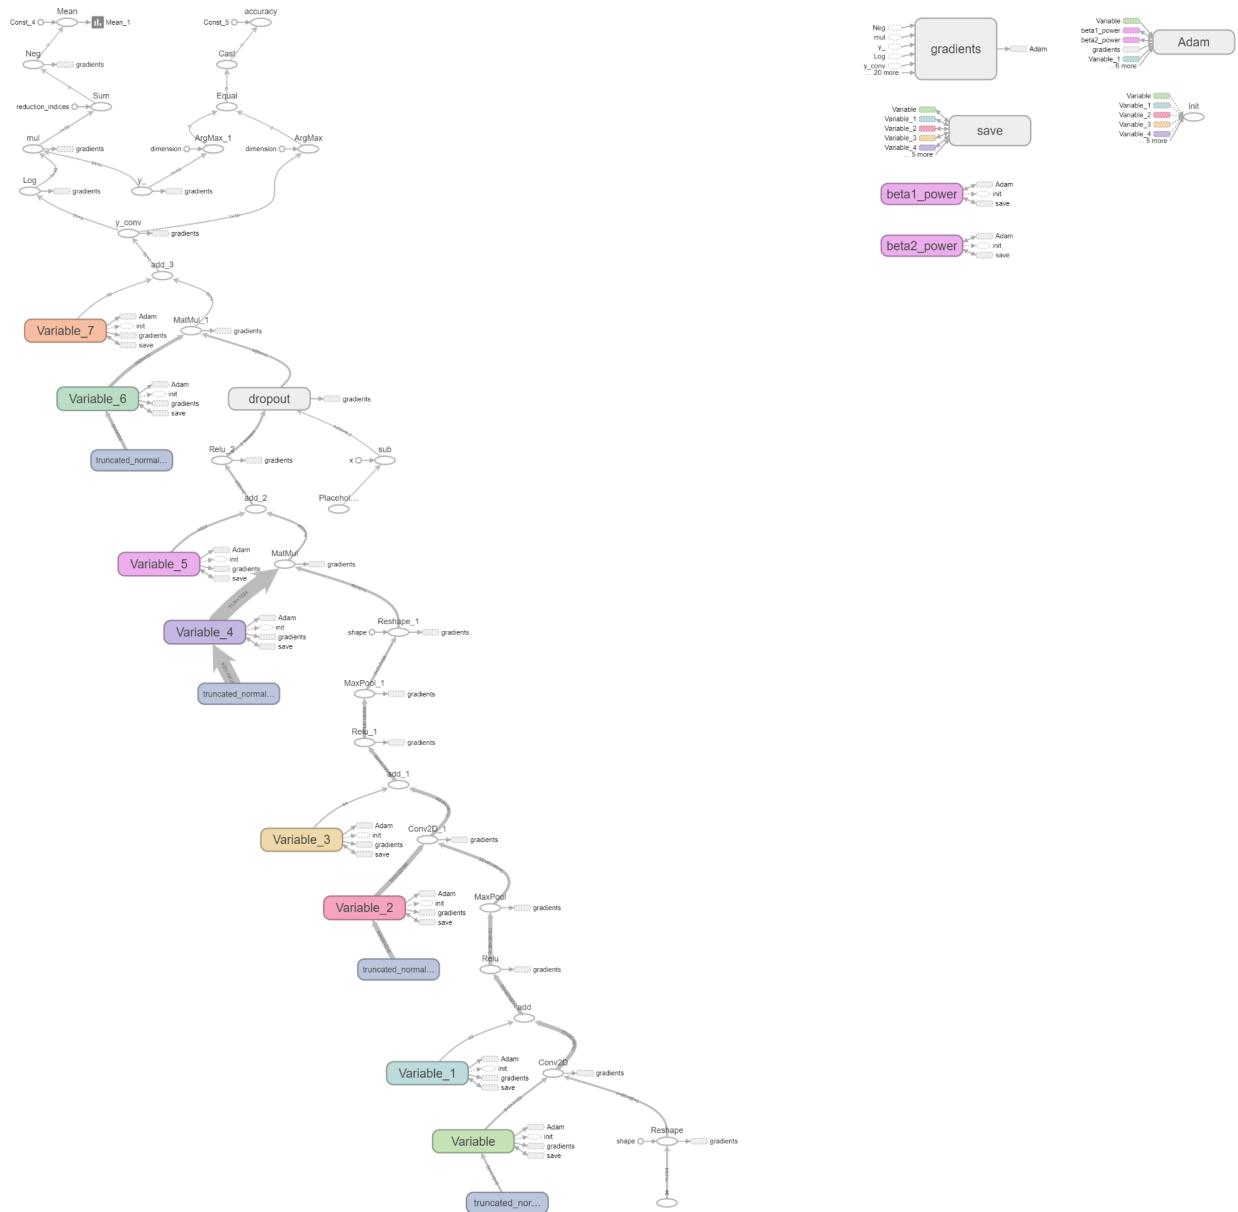


Figure 2\_2: graphs

The figure of projector:

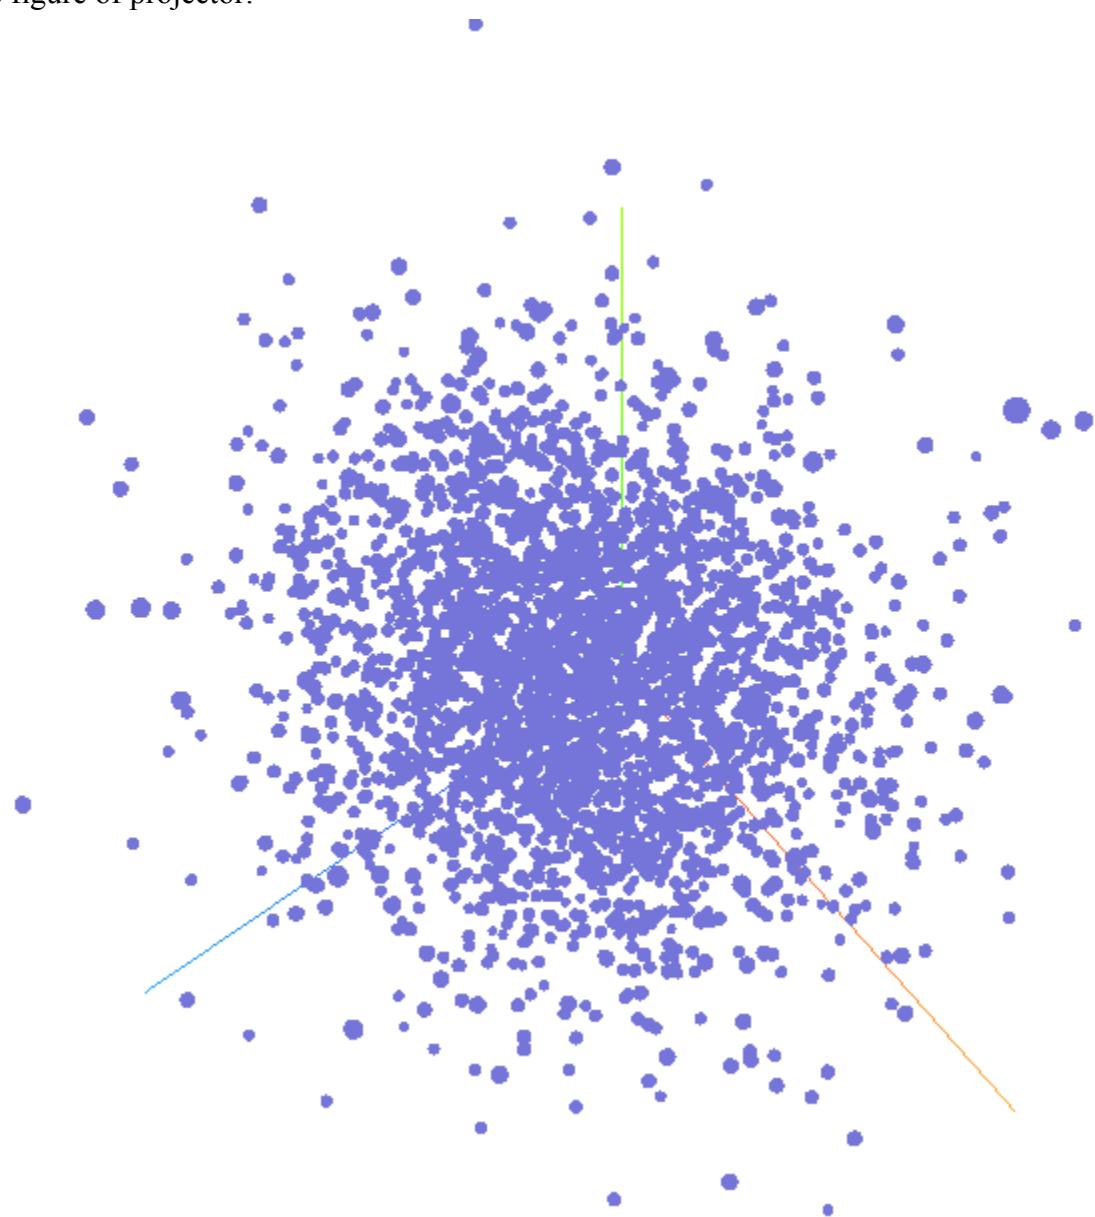


Figure 2\_3: projectors

2.b

Write function variable\_summaries for plotting variables:

```

def variable_summaries(var):
    """
    Attach a lot of summaries to a Tensor (for TensorBoard visualization).
    """

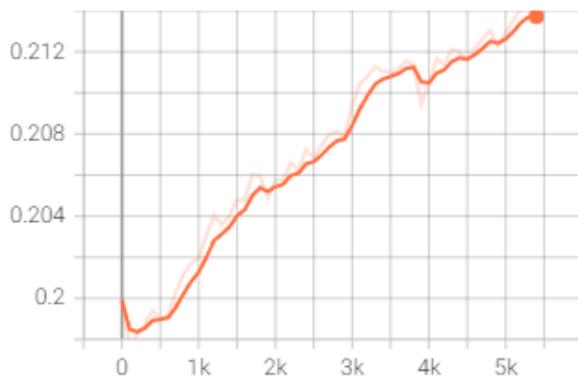
    with tf.name_scope('summaries'):
        mean = tf.reduce_mean(var)
        tf.summary.scalar('mean', mean)
        with tf.name_scope('stddev'):
            stddev = tf.sqrt(tf.reduce_mean(tf.square(var - mean)))
        tf.summary.scalar('stddev', stddev)
        tf.summary.scalar('max', tf.reduce_max(var))
        tf.summary.scalar('min', tf.reduce_min(var))
    tf.summary.histogram('histogram', var)

```

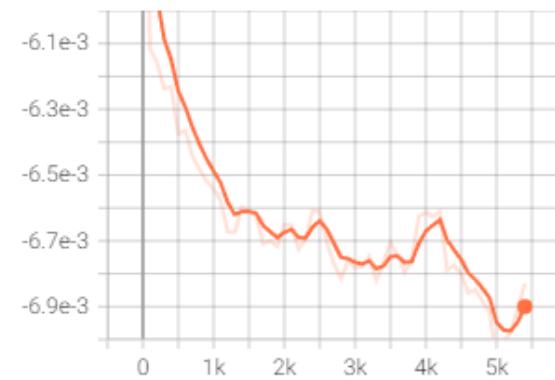
To use this function, we run it on W\_conv1, b\_conv1, h\_conv1, h\_pool1, W\_conv2, b\_conv2, h\_conv2, h\_pool2, W\_fc1, b\_fc1, h\_pool2\_flat, h\_fc1, keep\_prob, h\_fc1\_drop, W\_fc2, b\_fc2, and y\_conv.

**Here** are the results:

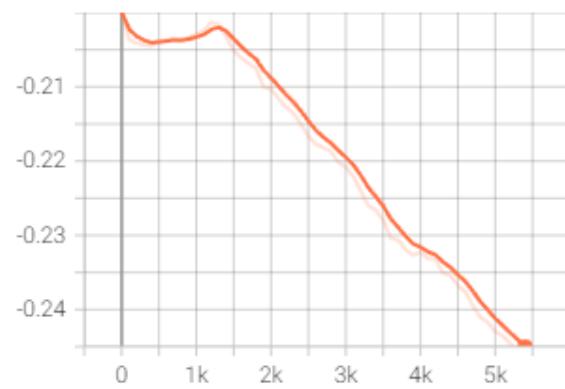
summaries/max\_1  
tag: summaries/max\_1



summaries/mean\_1  
tag: summaries/mean\_1



summaries/min\_1  
tag: summaries/min\_1



summaries/stddev\_1  
tag: summaries/stddev\_1

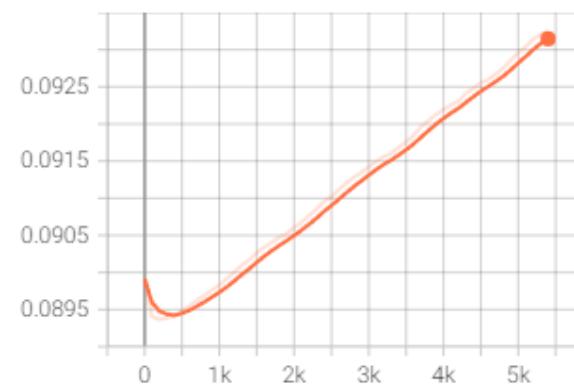


Figure 2\_4: W\_conv1

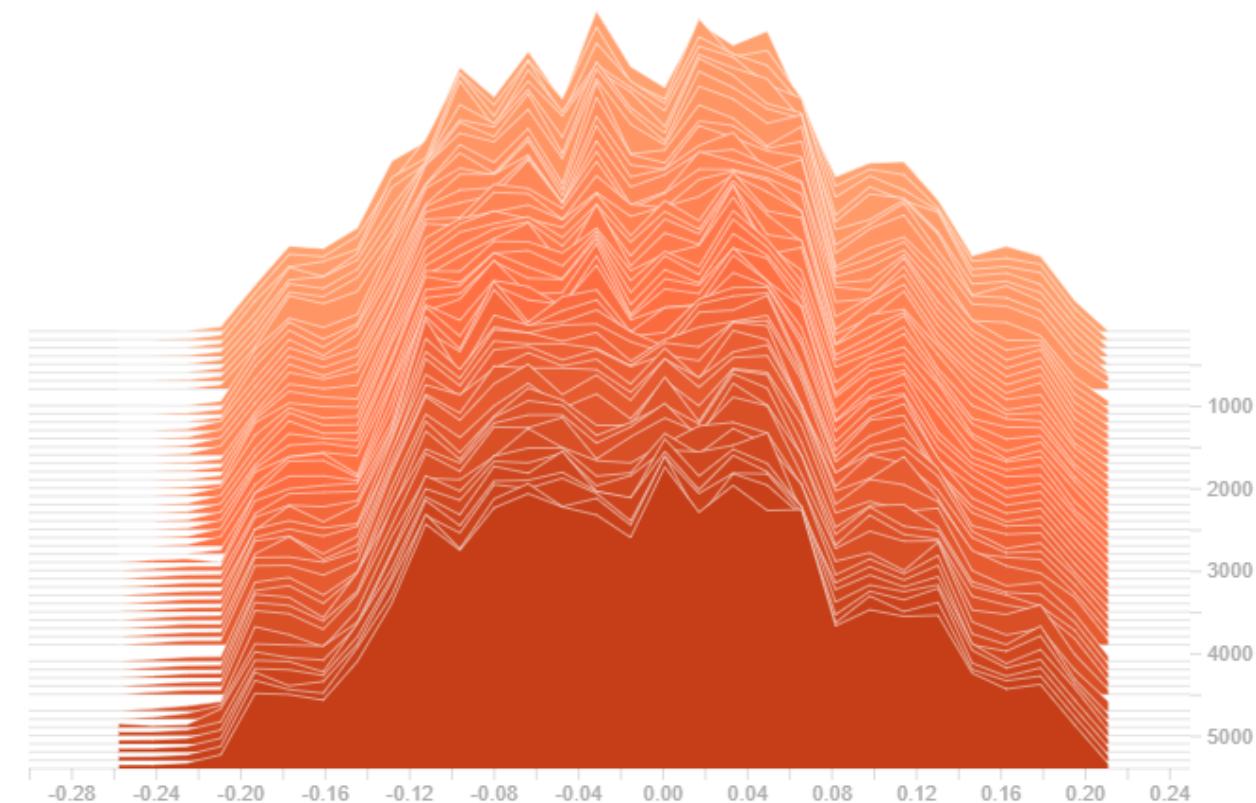


Figure 2\_5: W\_conv1\_histogram

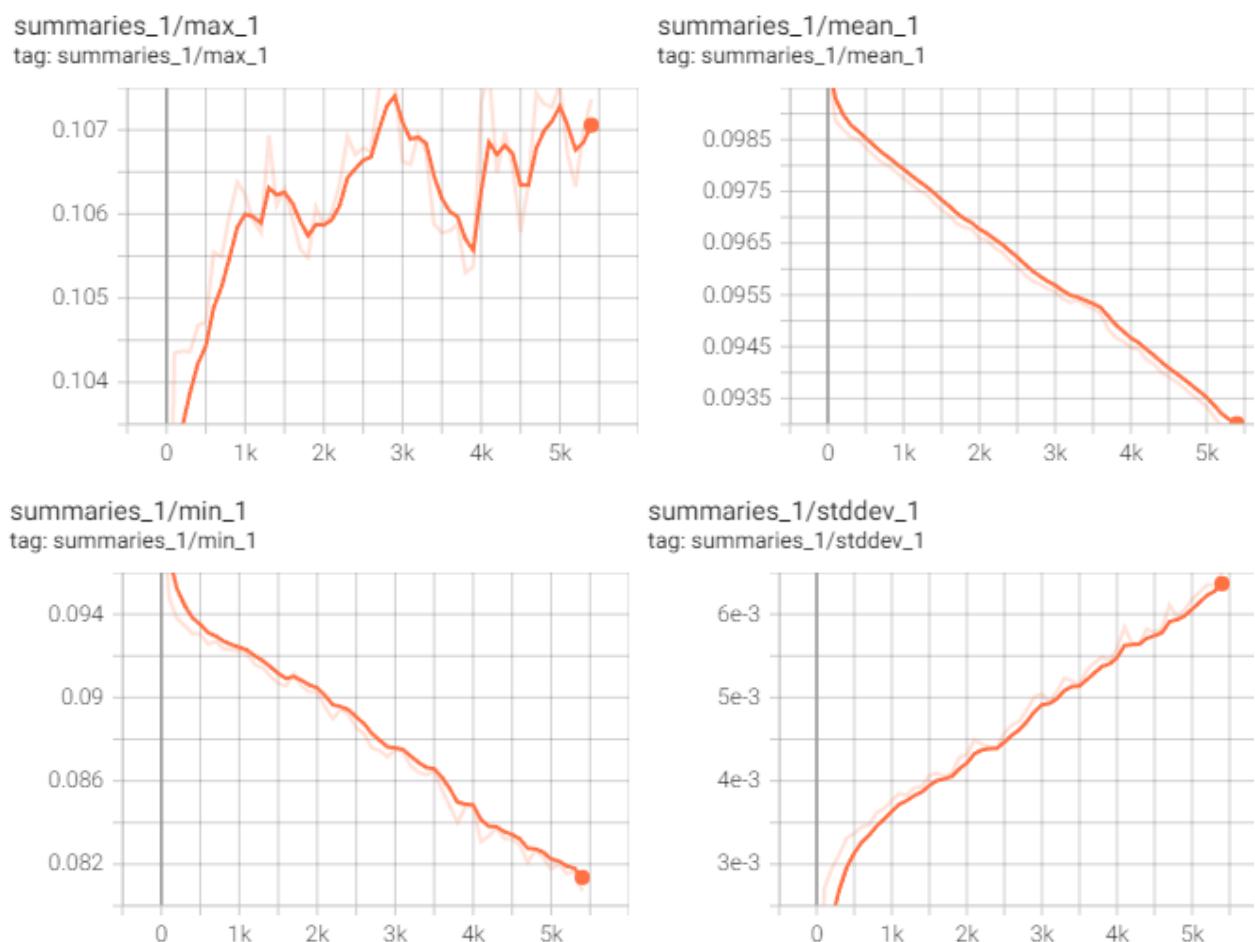


Figure 2\_6: b\_conv1

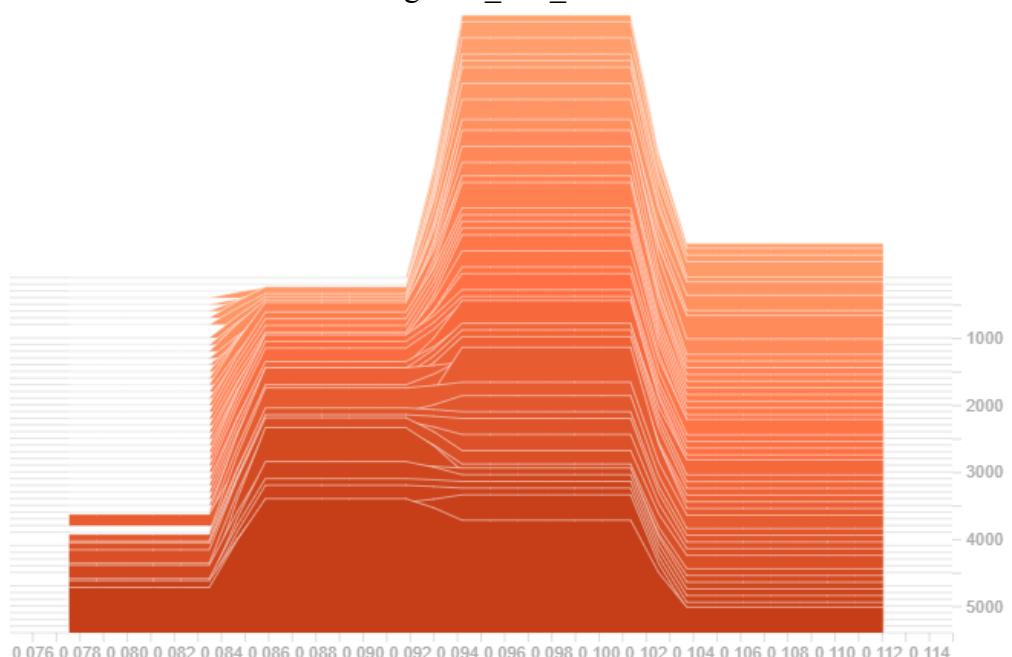


Figure 2\_7: b\_conv1\_histogram

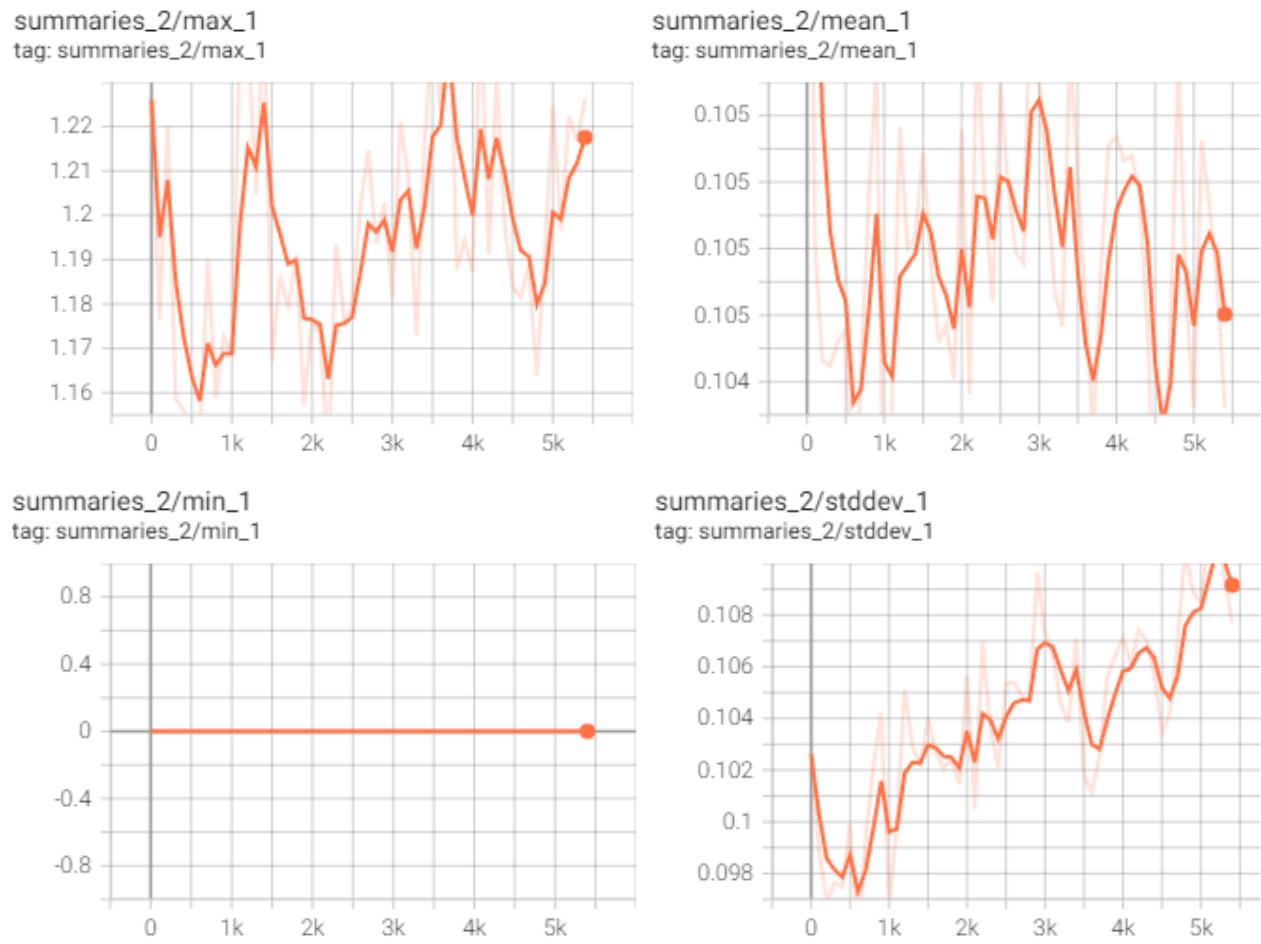


Figure 2\_8: h\_conv1

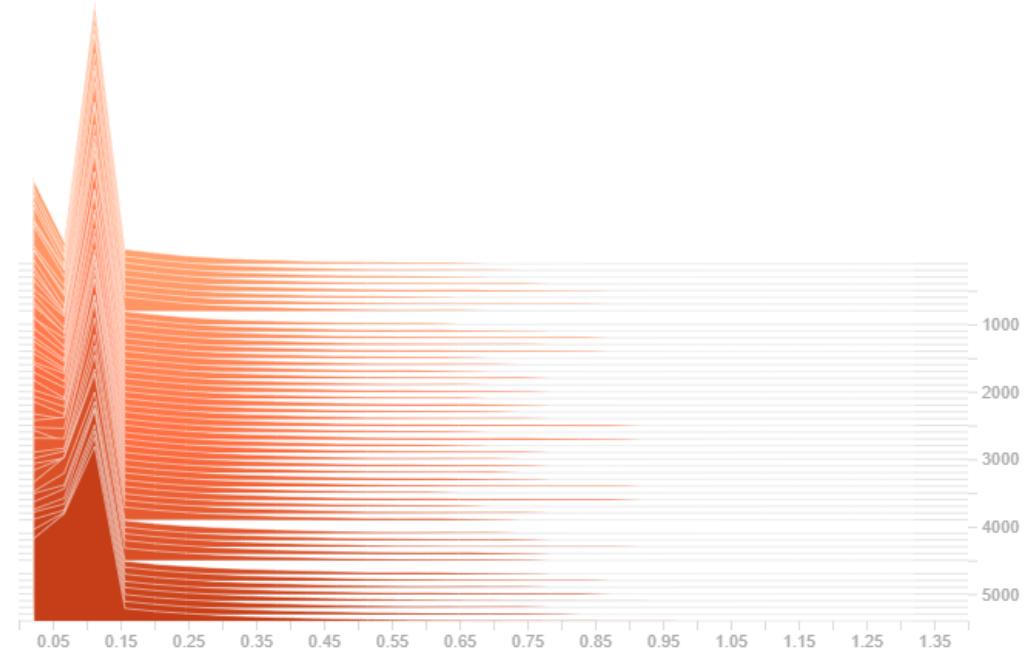


Figure 2\_9: h\_conv1\_histogram

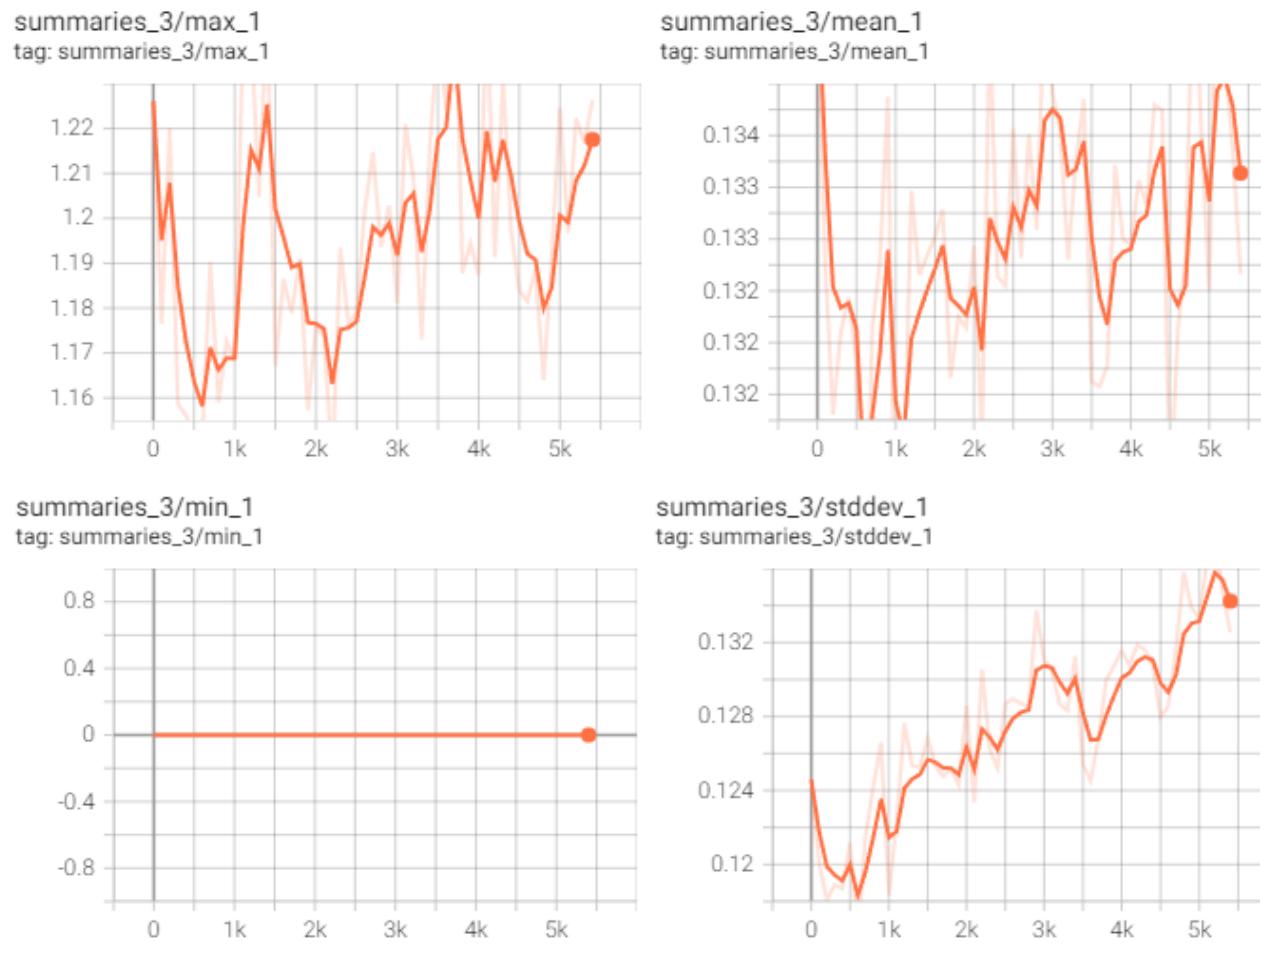


Figure 2\_10: h\_pool1

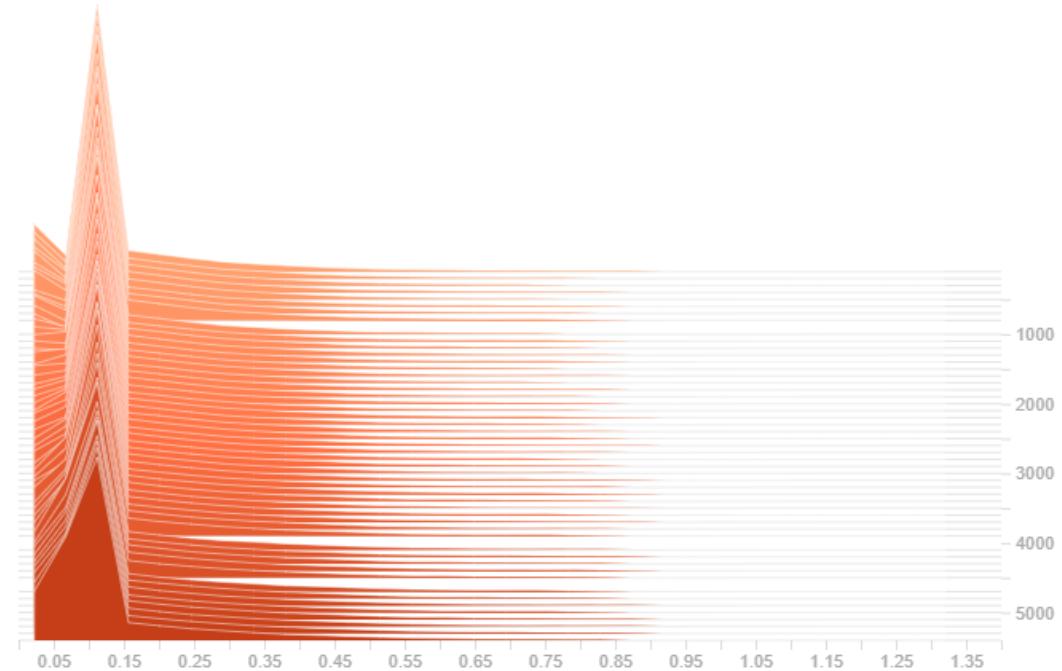
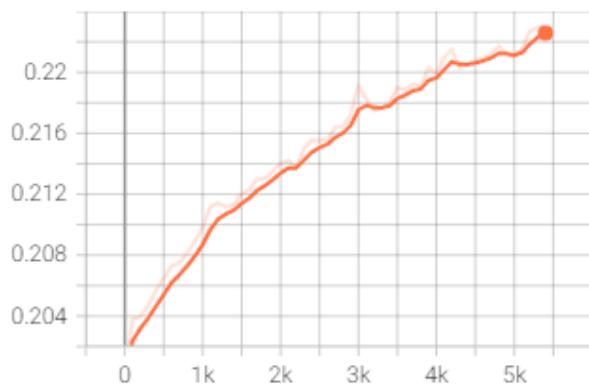
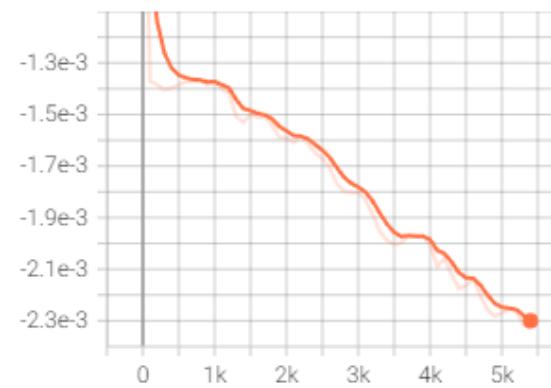


Figure 2\_11: h\_pool1\_histogram

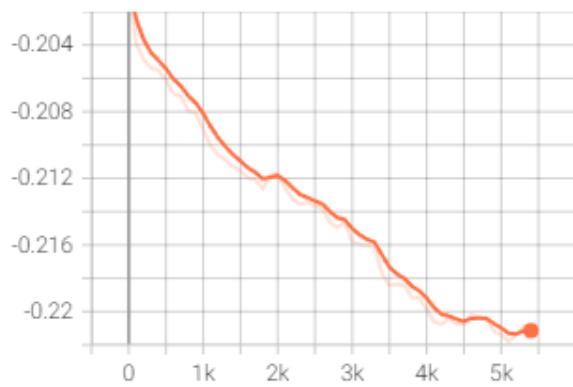
summaries\_4/max\_1  
tag: summaries\_4/max\_1



summaries\_4/mean\_1  
tag: summaries\_4/mean\_1



summaries\_4/min\_1  
tag: summaries\_4/min\_1



summaries\_4/stddev\_1  
tag: summaries\_4/stddev\_1

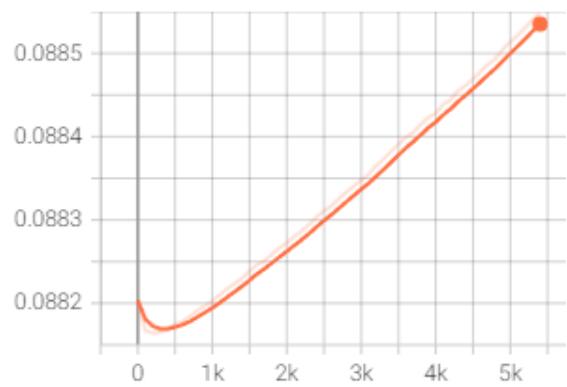


Figure 2\_12: W\_conv2

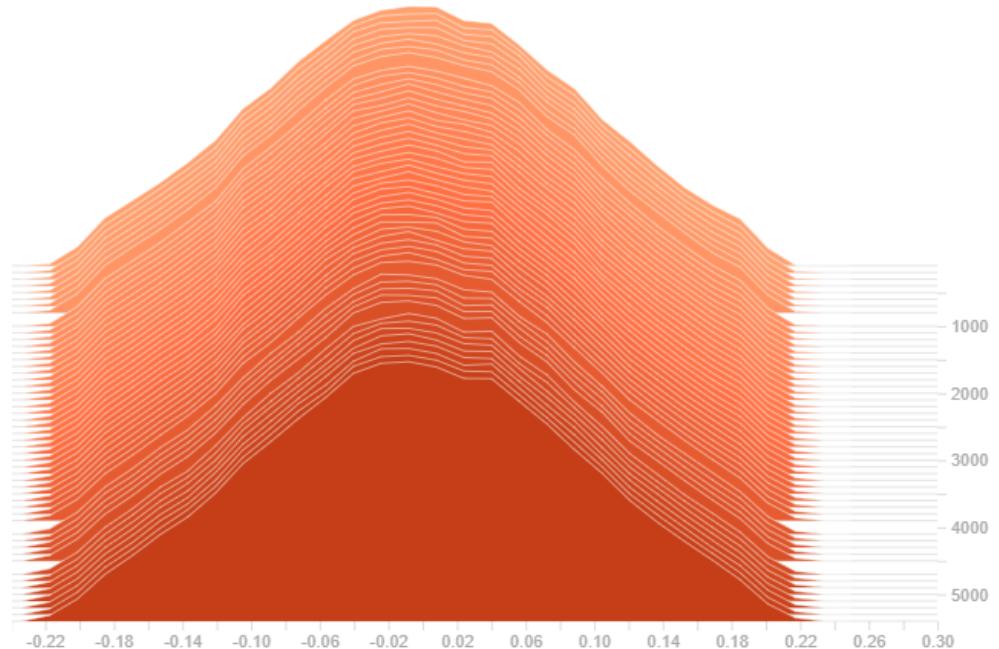


Figure 2\_13: W\_conv2\_histogram

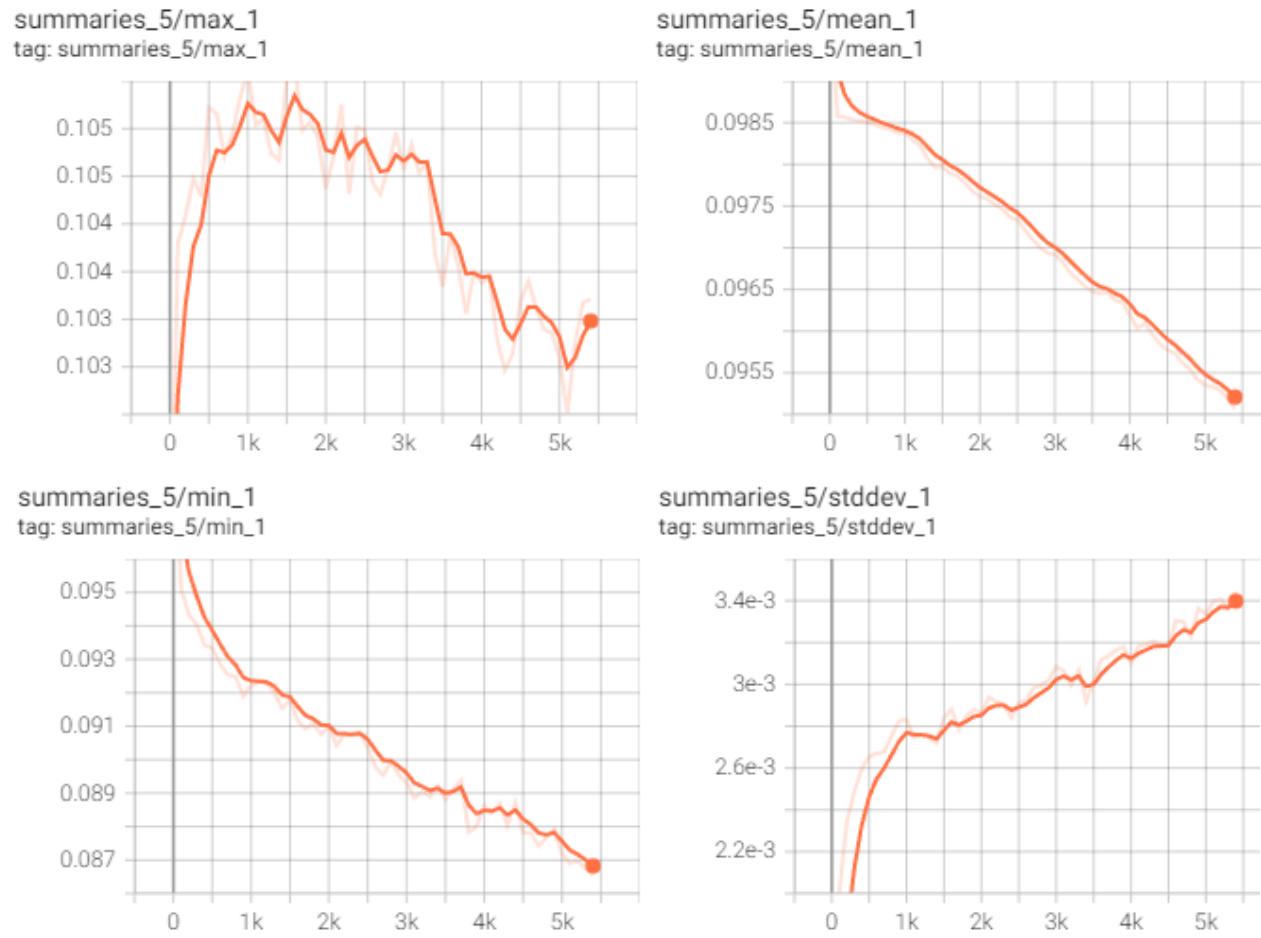


Figure 2\_14: b\_conv2

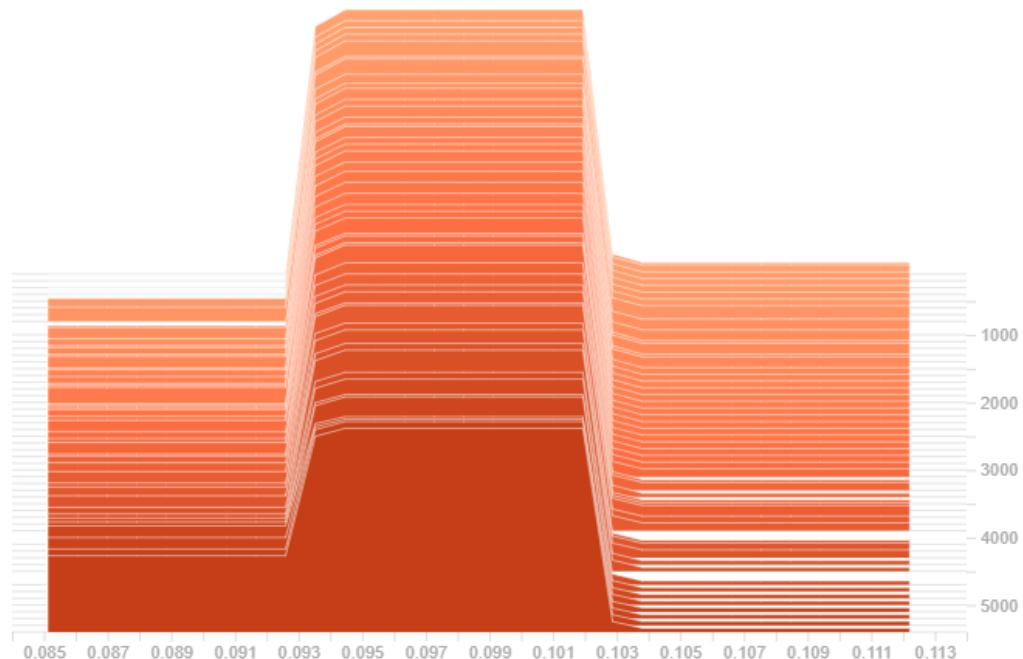


Figure 2\_15: b\_conv2\_histogram

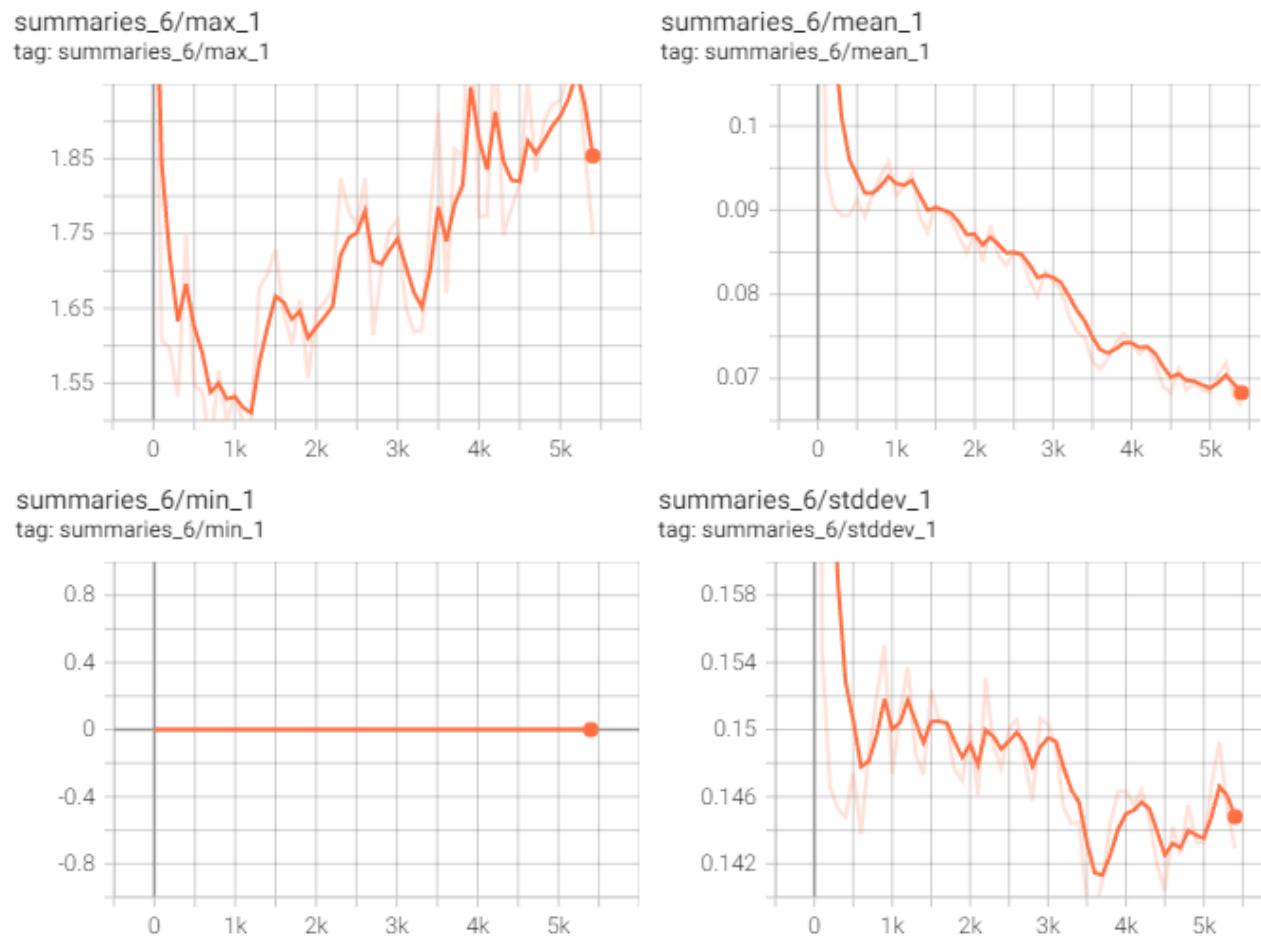


Figure 2\_16: `h_conv2`

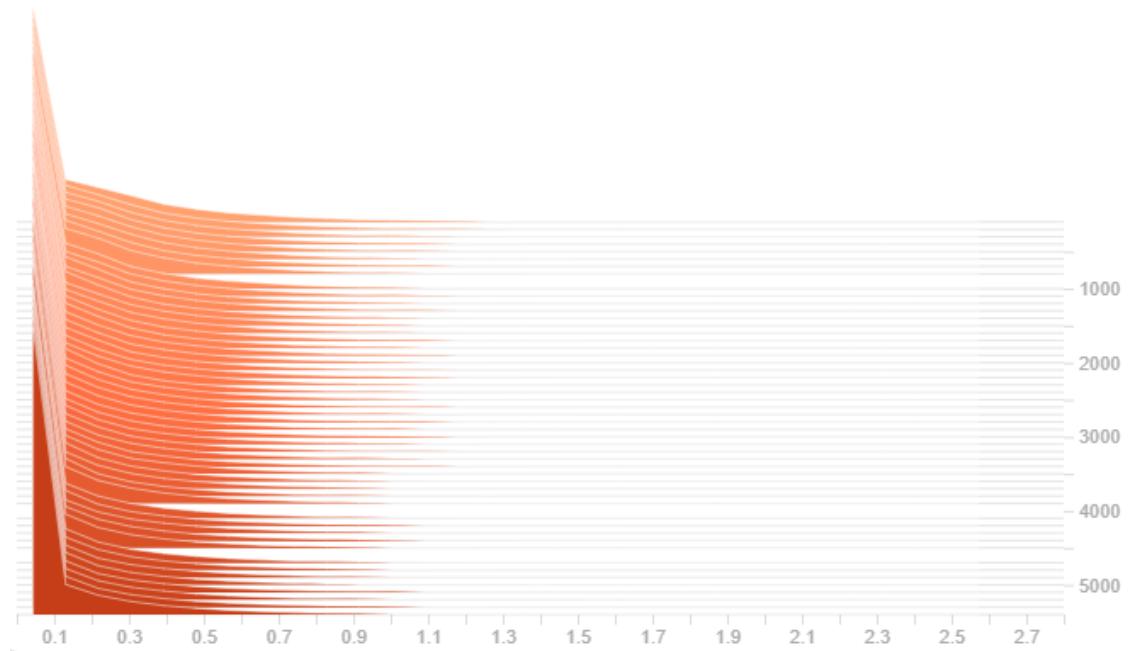


Figure 2\_17: `h_conv2_histogram`

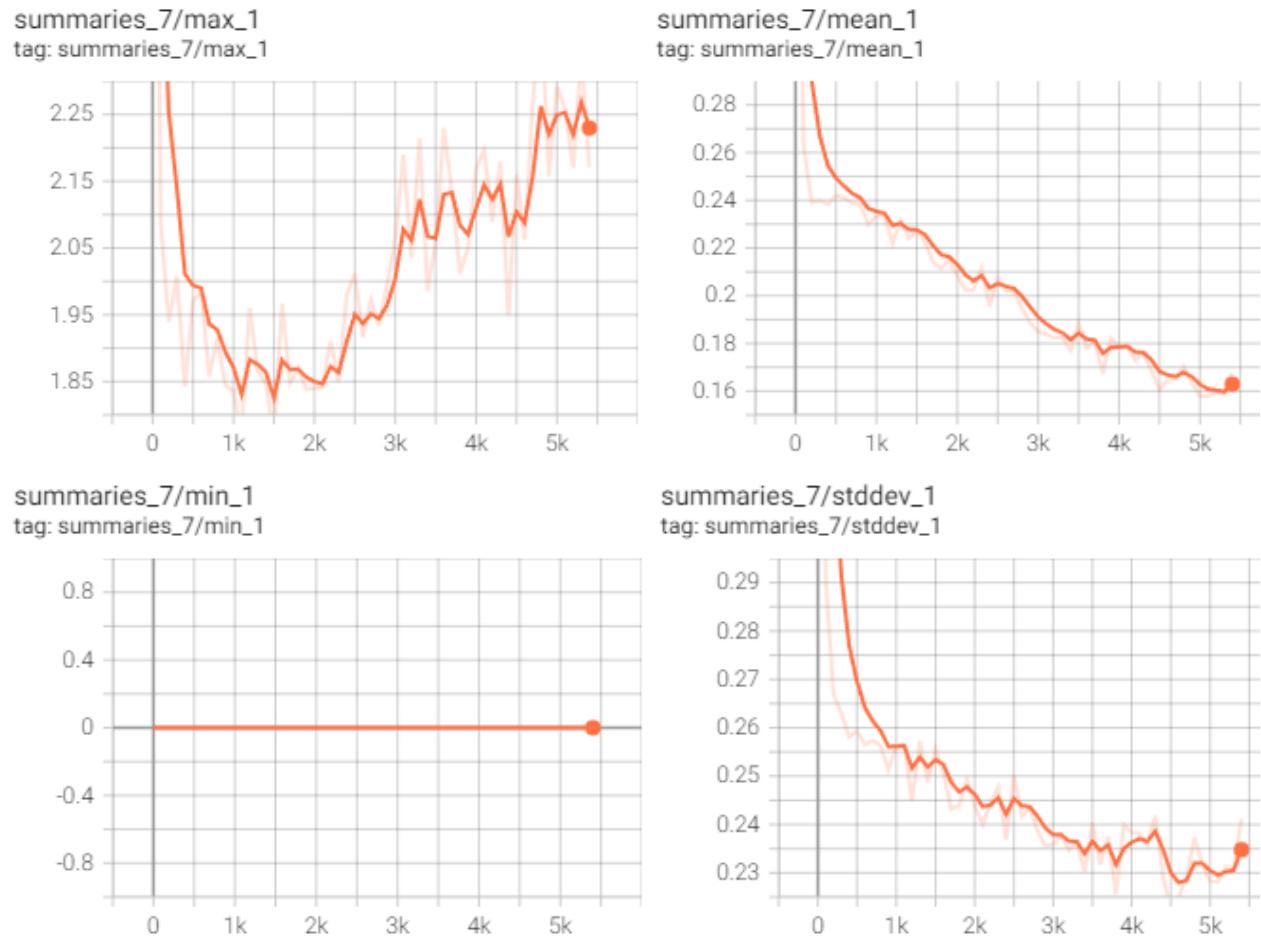


Figure 2\_18: h\_pool2

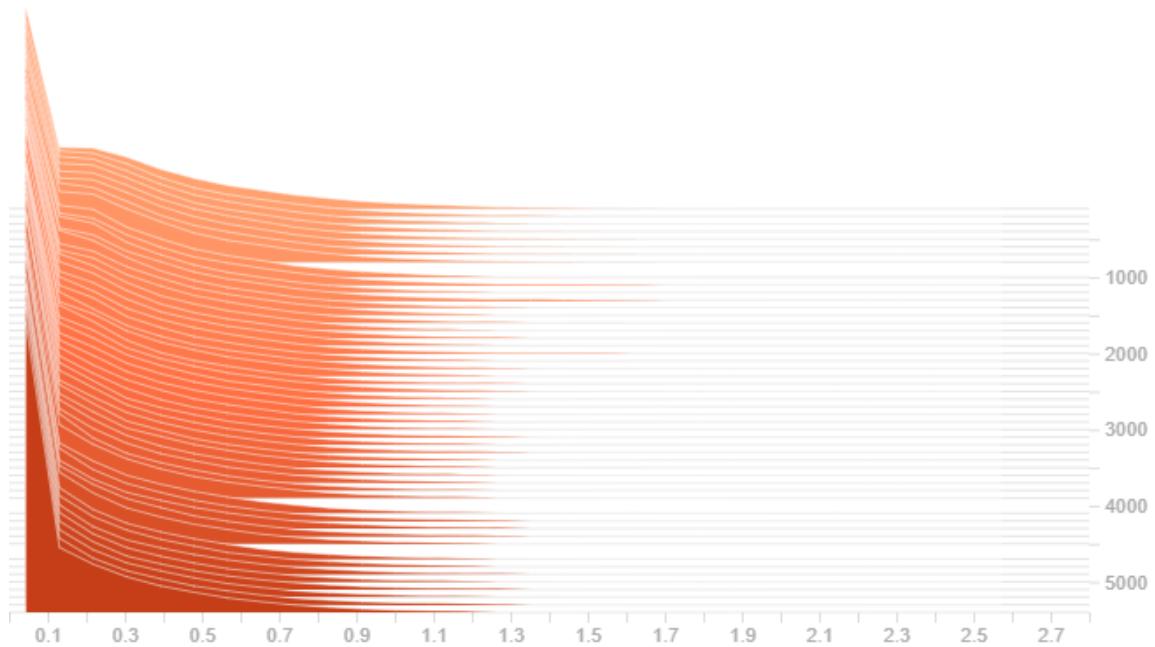


Figure 2\_19: h\_pool2\_histogram

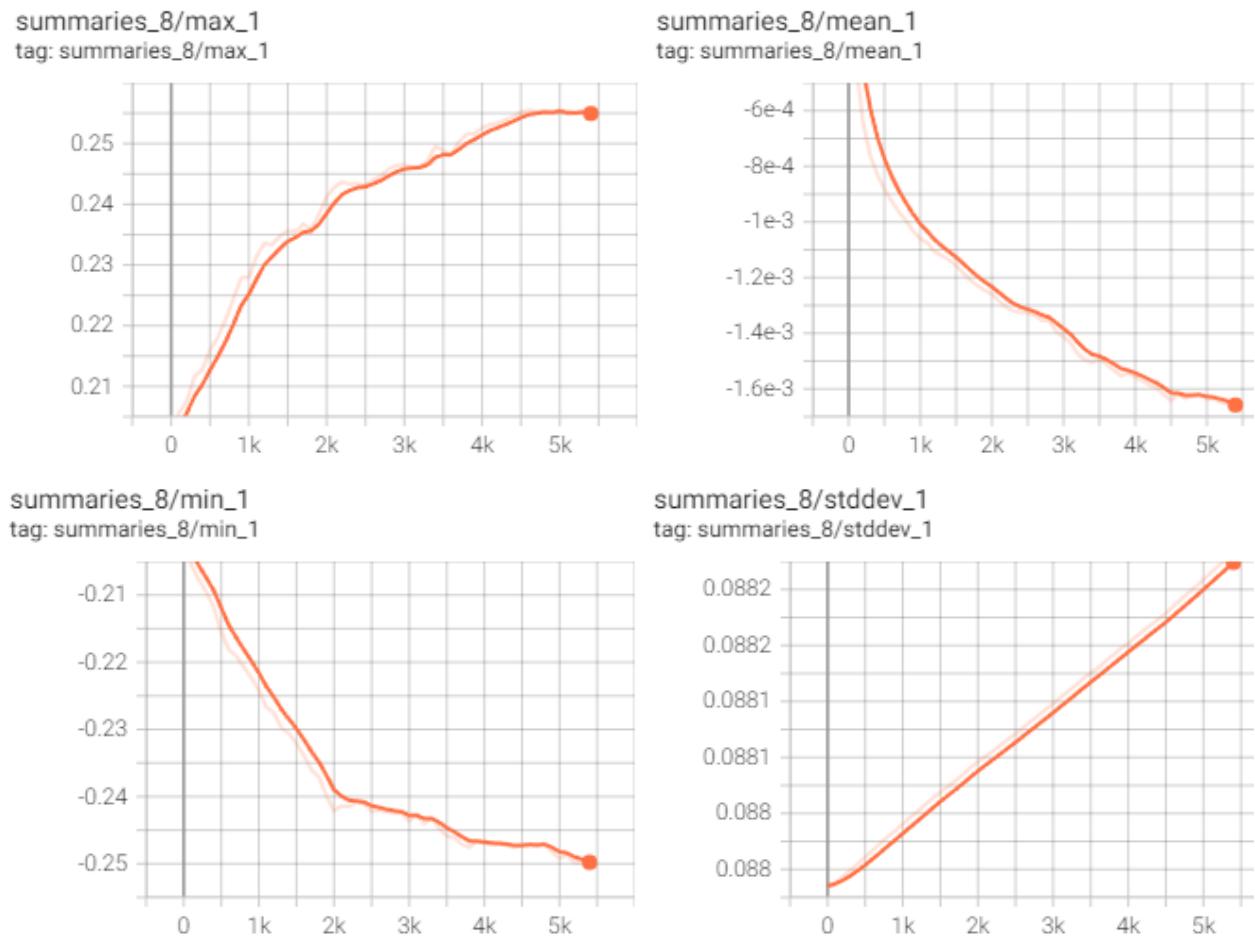


Figure 2\_20: W\_fc1

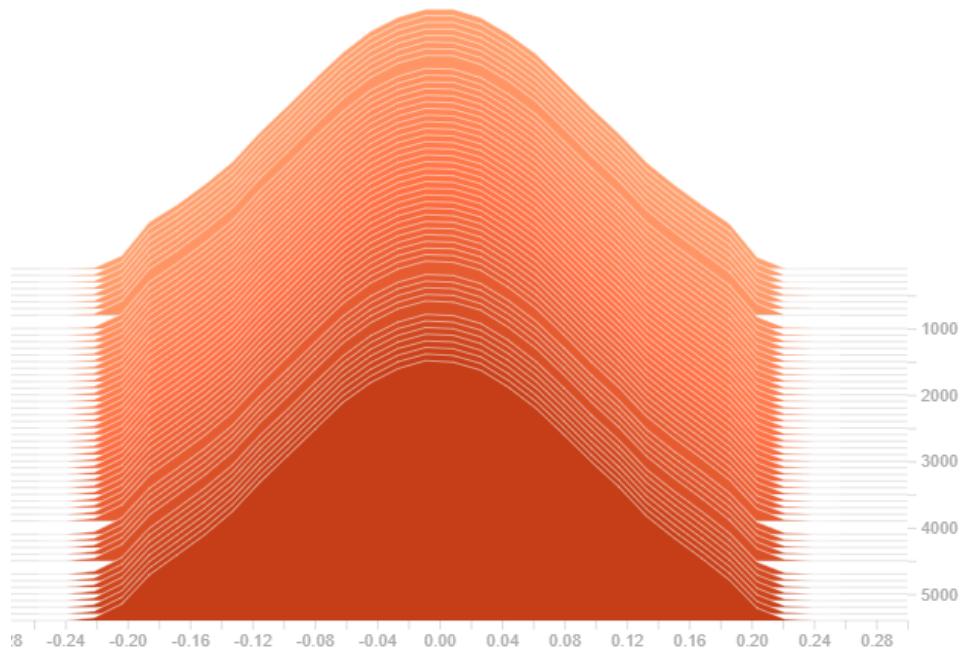


Figure 2\_21: W\_fc1\_histogram

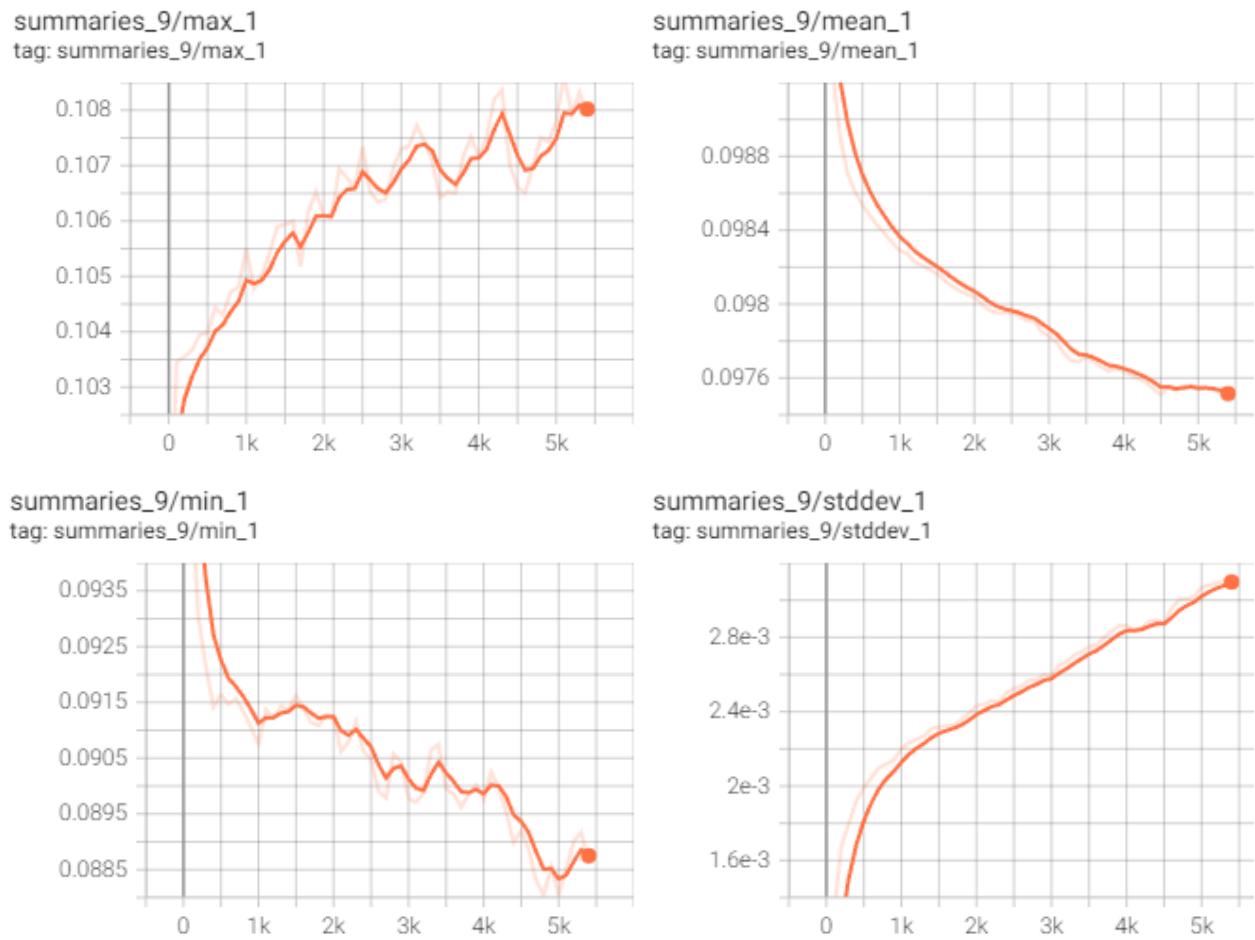


Figure 2\_22: b\_fc1

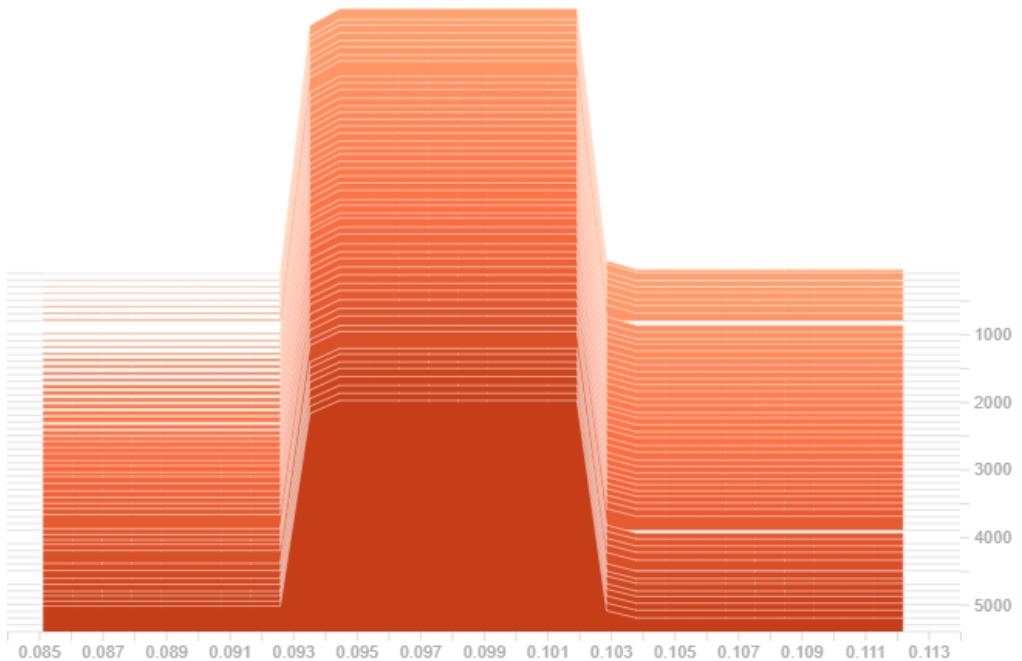


Figure 2\_23: b\_fc1\_histogram

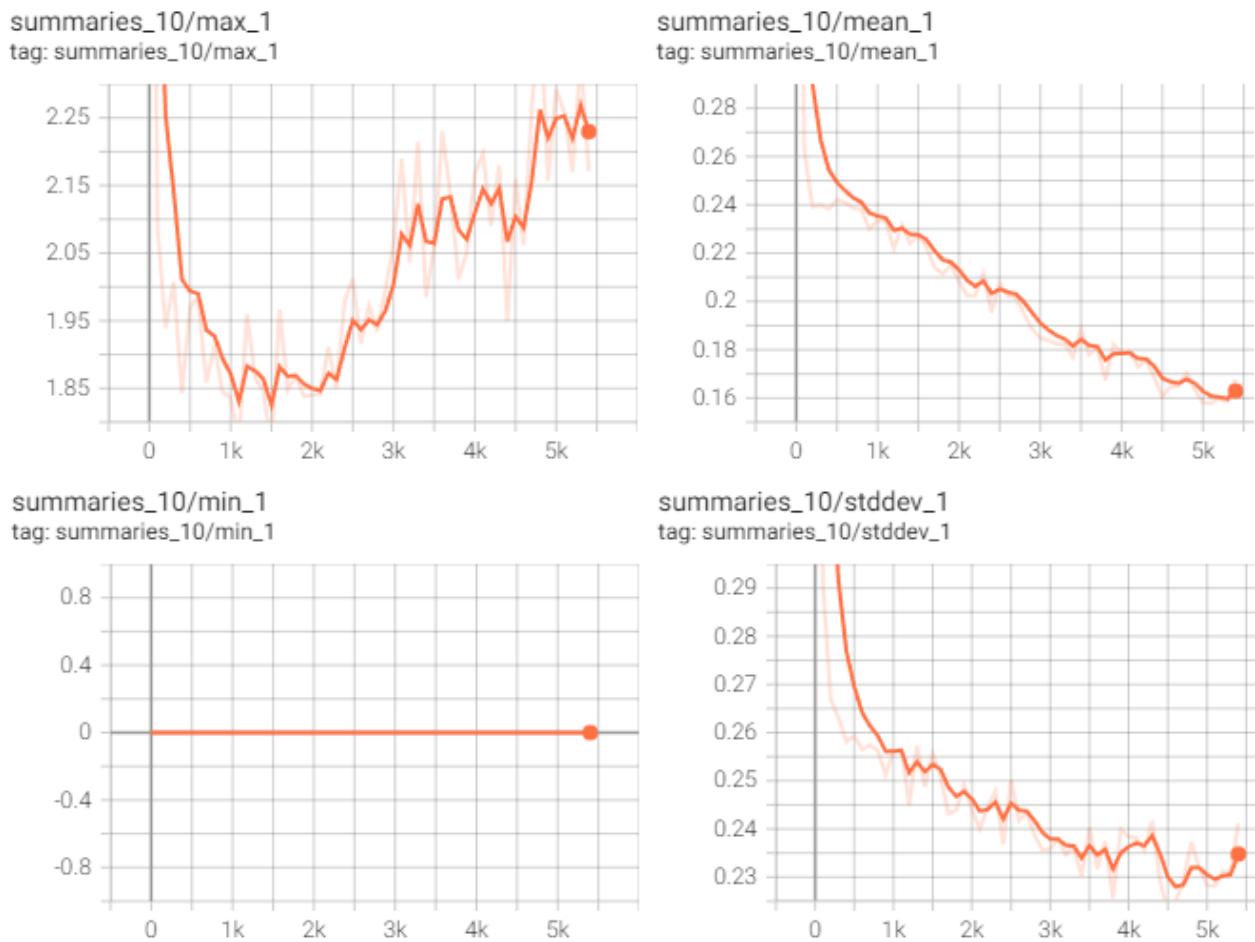


Figure 2\_24: `h_pool2_flat`

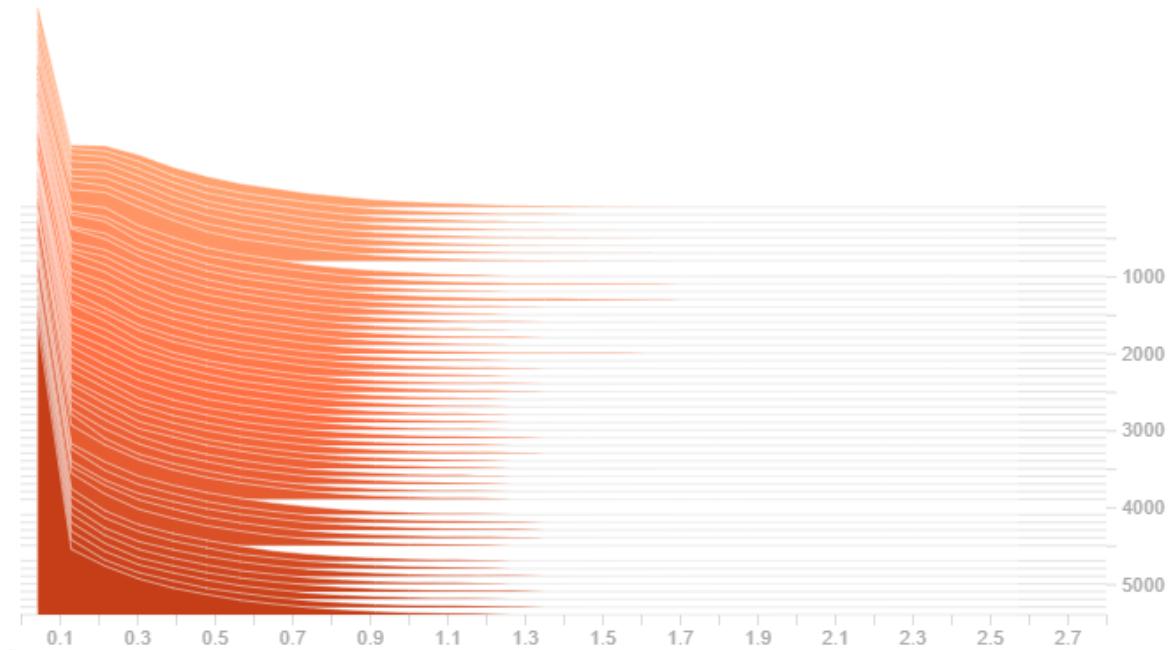
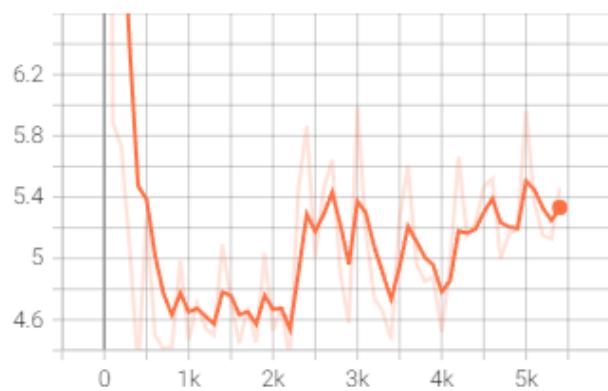
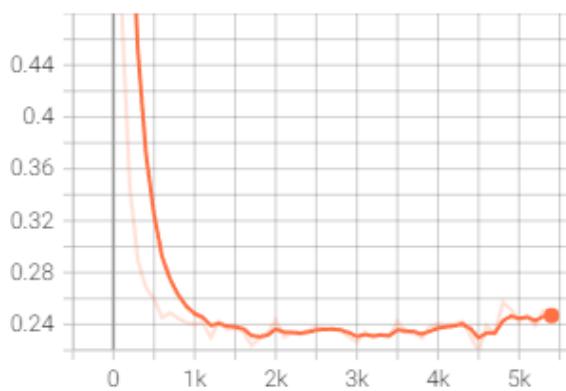


Figure 2\_25: `h_pool2_flat_histogram`

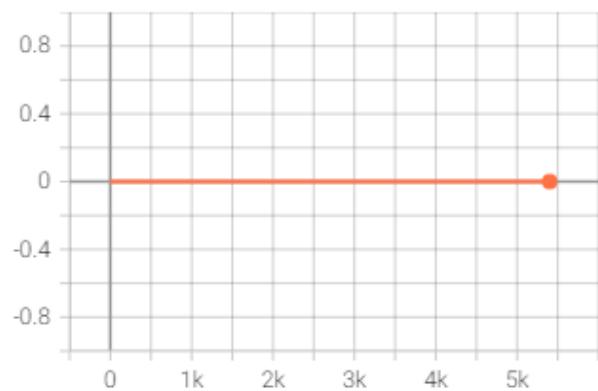
summaries\_11/max\_1  
tag: summaries\_11/max\_1



summaries\_11/mean\_1  
tag: summaries\_11/mean\_1



summaries\_11/min\_1  
tag: summaries\_11/min\_1



summaries\_11/stddev\_1  
tag: summaries\_11/stddev\_1

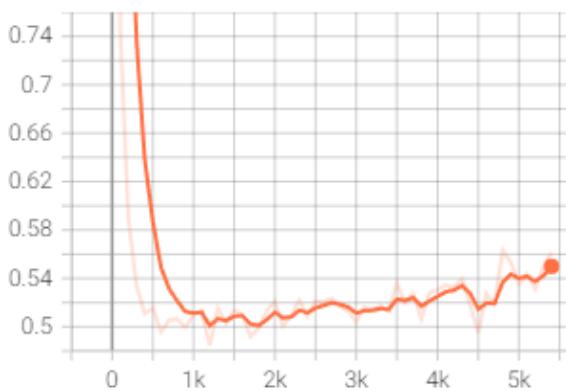


Figure 2\_26: h\_fc1

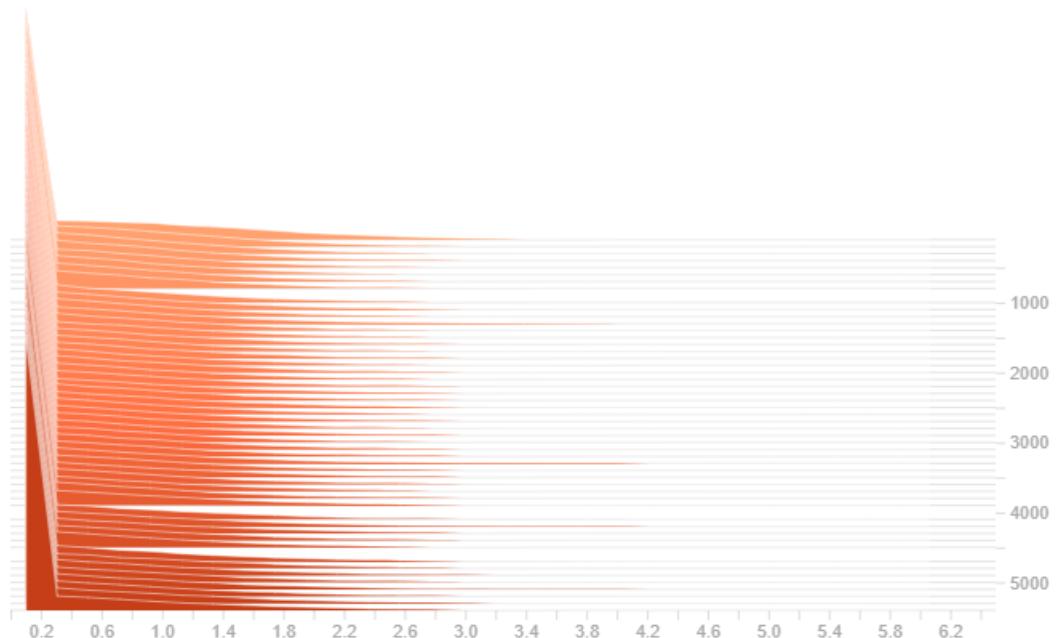
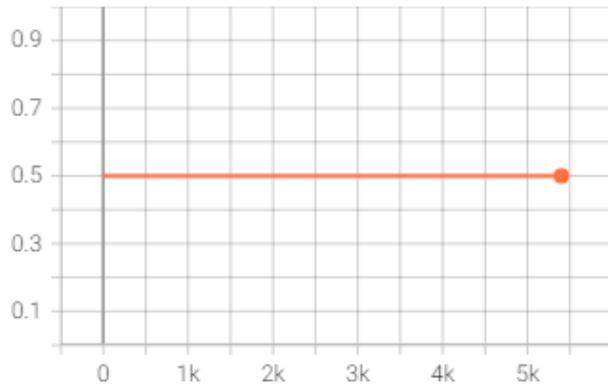
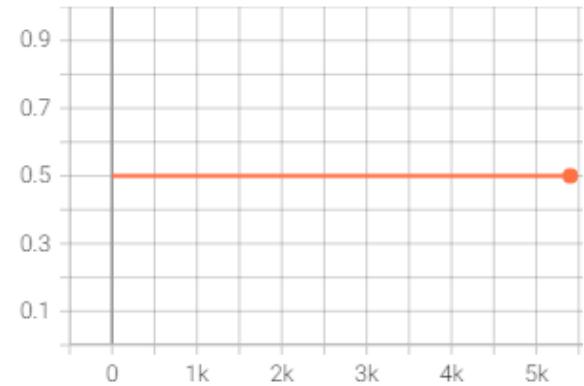


Figure 2\_27: h\_fc1\_histogram

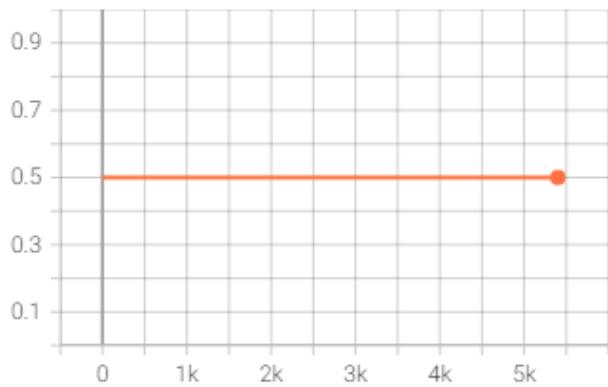
summaries\_12/max\_1  
tag: summaries\_12/max\_1



summaries\_12/mean\_1  
tag: summaries\_12/mean\_1



summaries\_12/min\_1  
tag: summaries\_12/min\_1



summaries\_12/stddev\_1  
tag: summaries\_12/stddev\_1

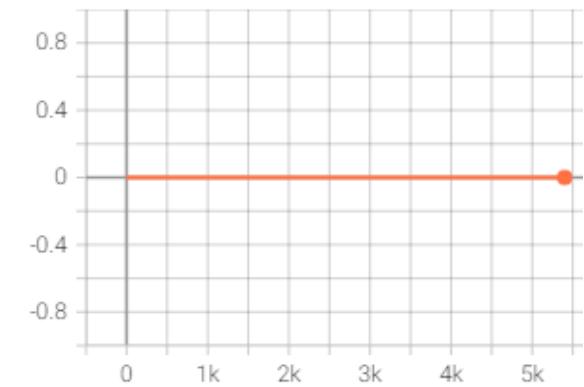


Figure 2\_28: keep\_prob

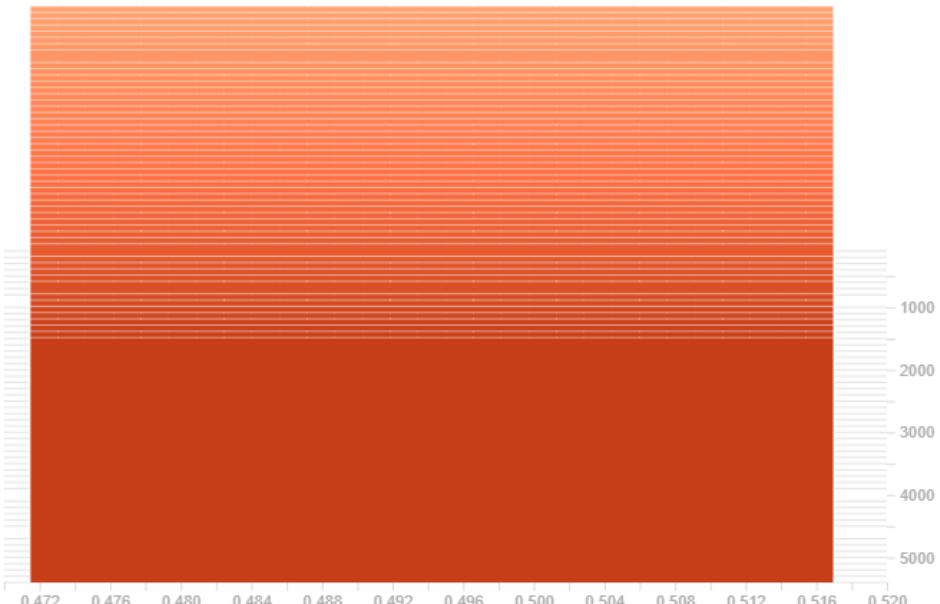
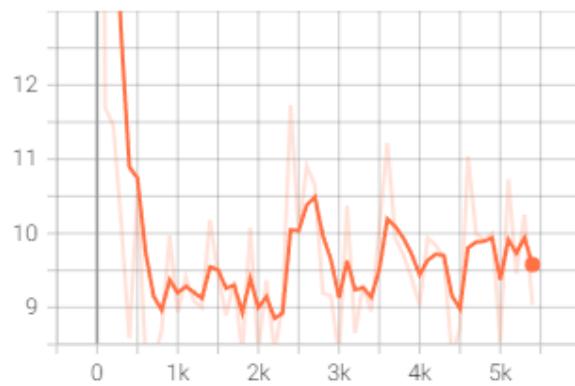


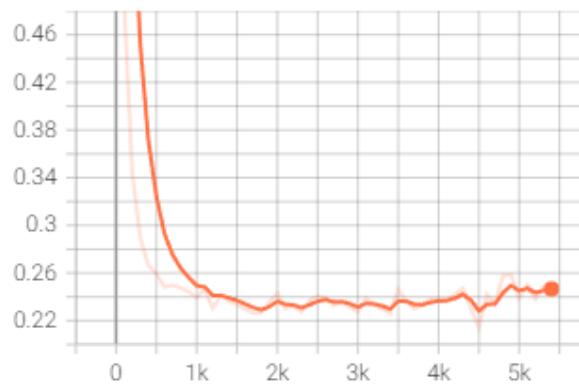
Figure 2\_29: keep\_prob\_histogram

summaries\_13/max\_1  
tag: summaries\_13/max\_1

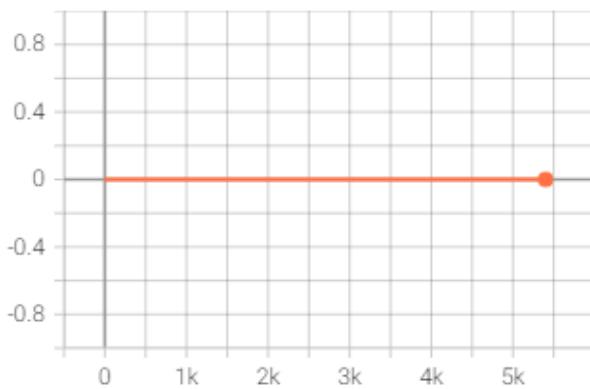


summaries\_13/mean\_1  
tag: summaries\_13/mean\_1

summaries\_13/mean\_1  
tag: summaries\_13/mean\_1



summaries\_13/min\_1  
tag: summaries\_13/min\_1



summaries\_13/stddev\_1  
tag: summaries\_13/stddev\_1

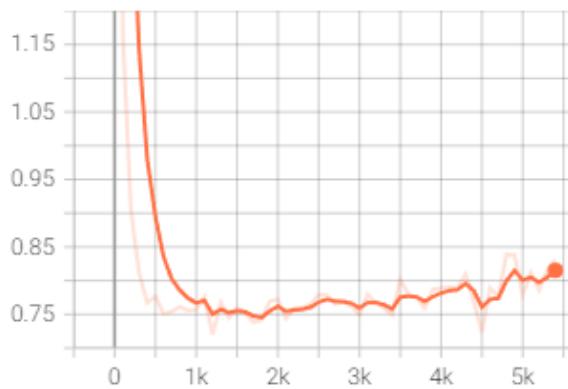


Figure 2\_30: h\_fc1\_drop

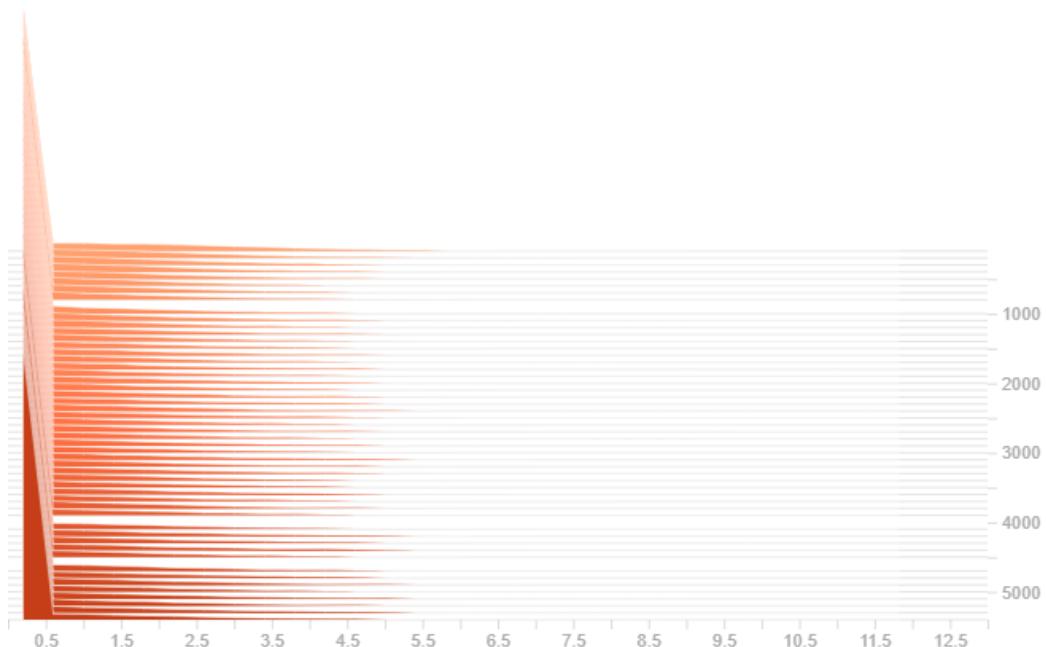


Figure 2\_31: h\_fc1\_drop\_histogram

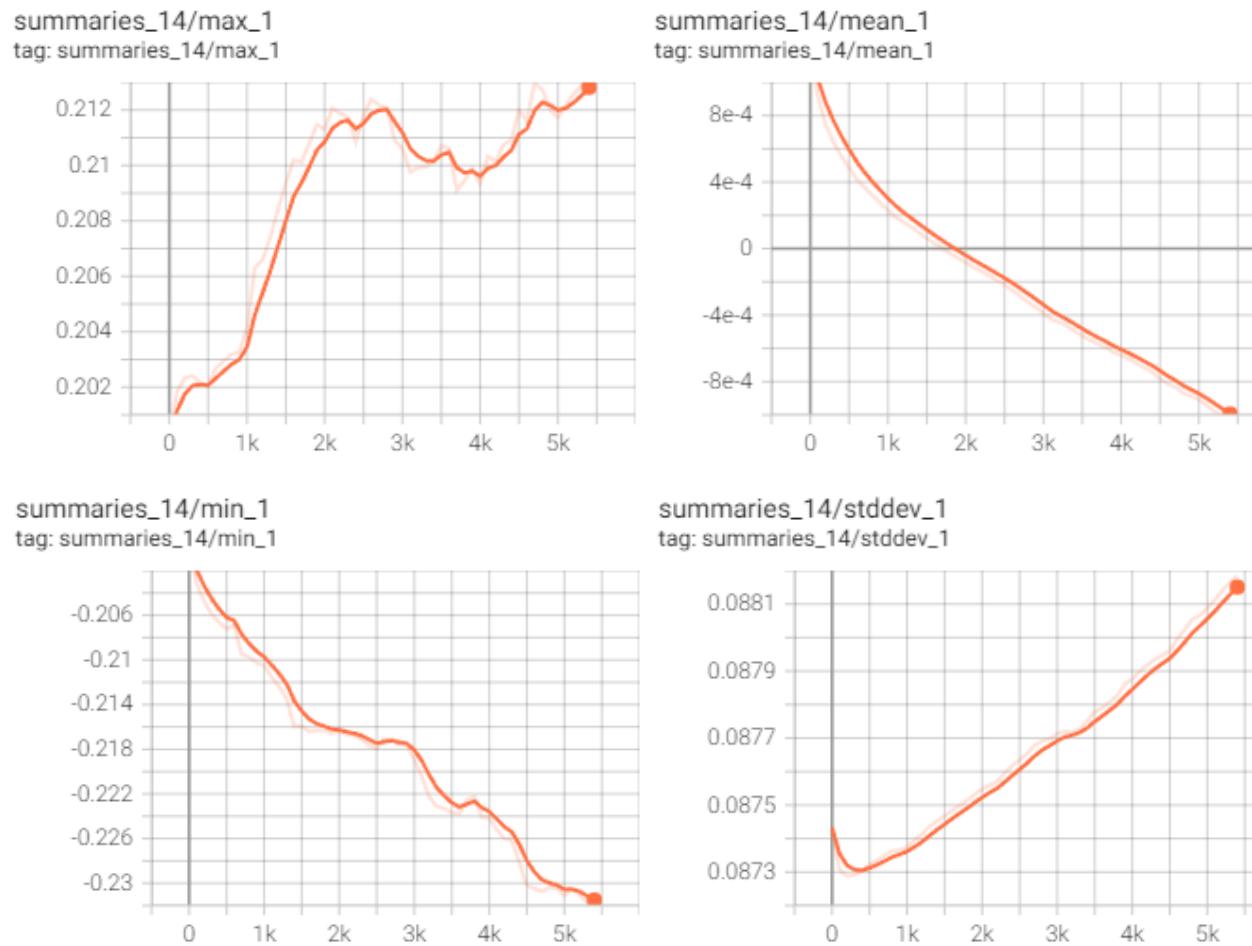


Figure 2\_32: W\_fc2

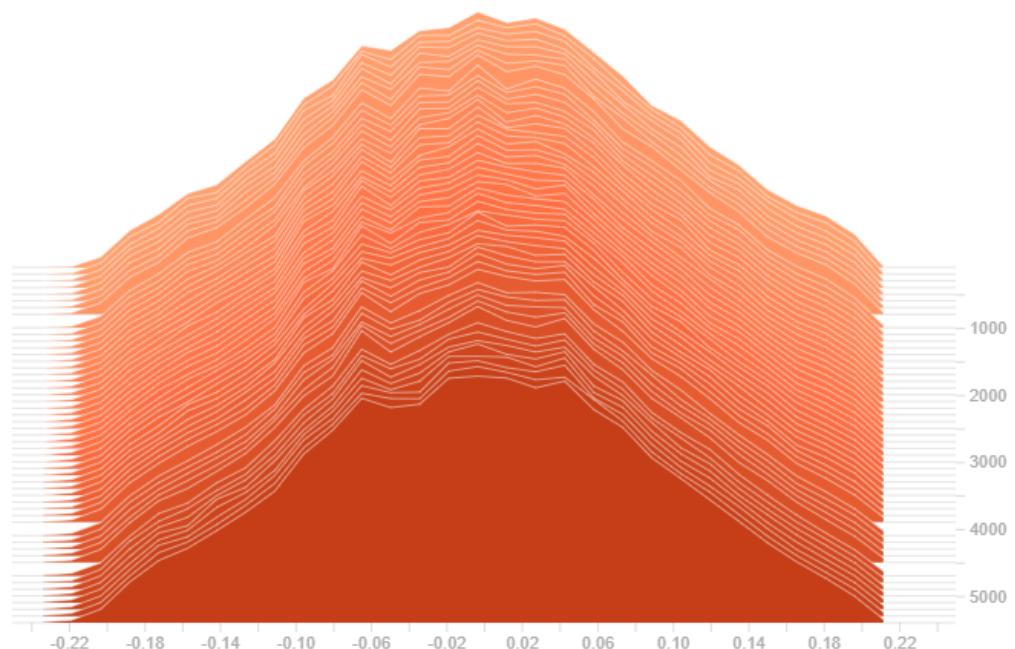


Figure 2\_33: W\_fc2\_histogram

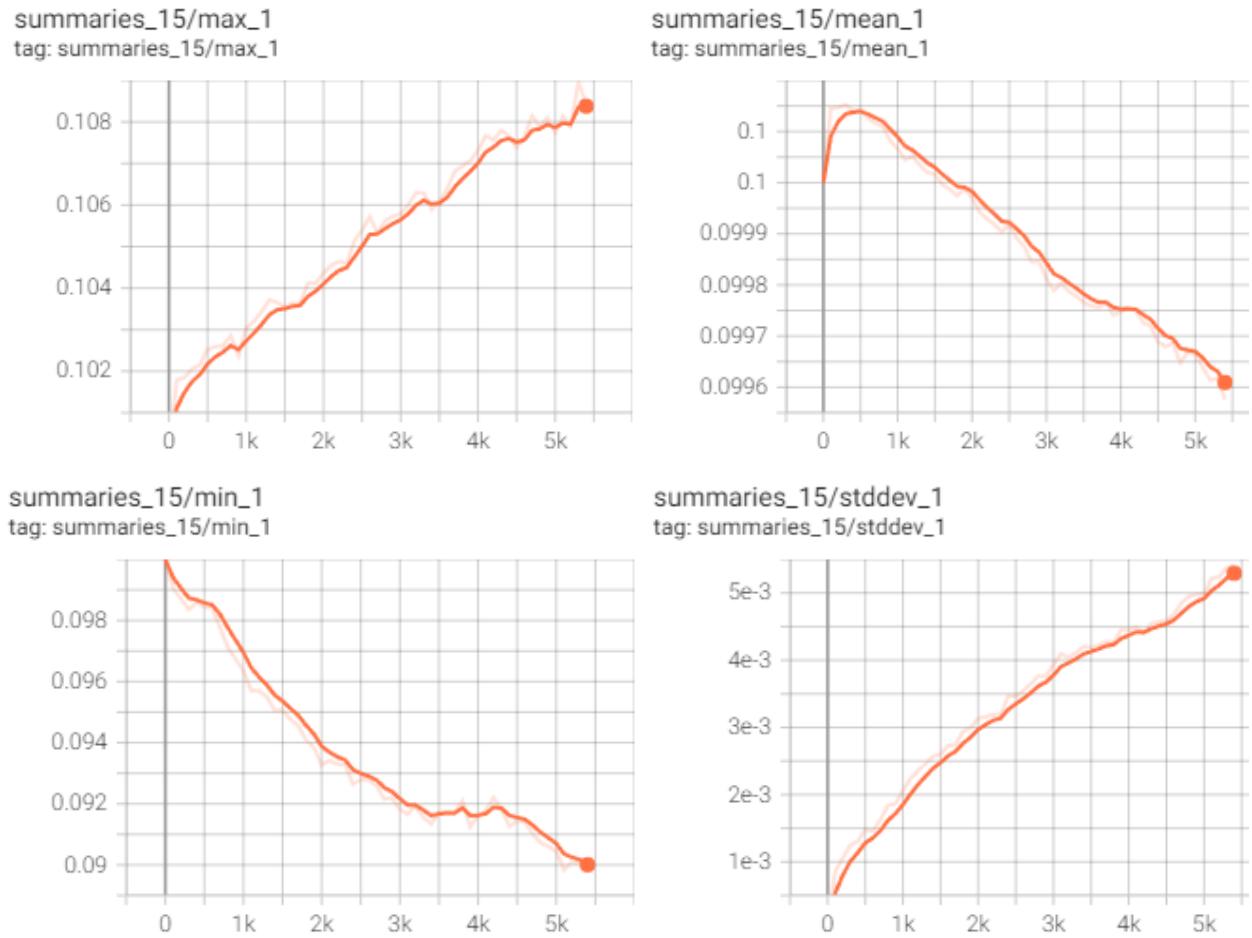


Figure 2\_34: b\_fc2

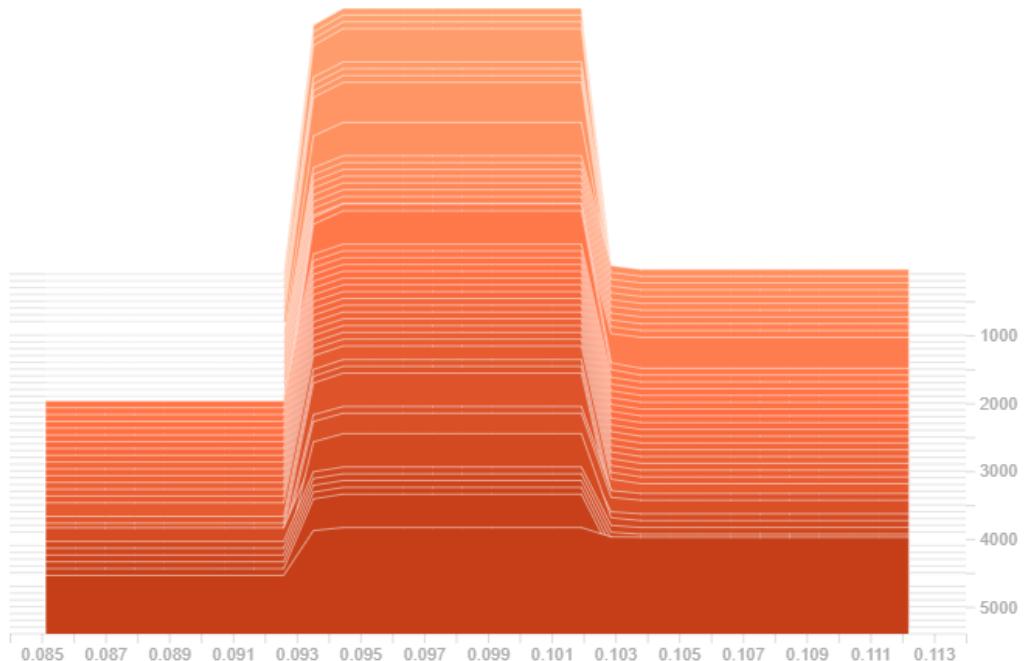


Figure 2\_35: b\_fc2\_histogram

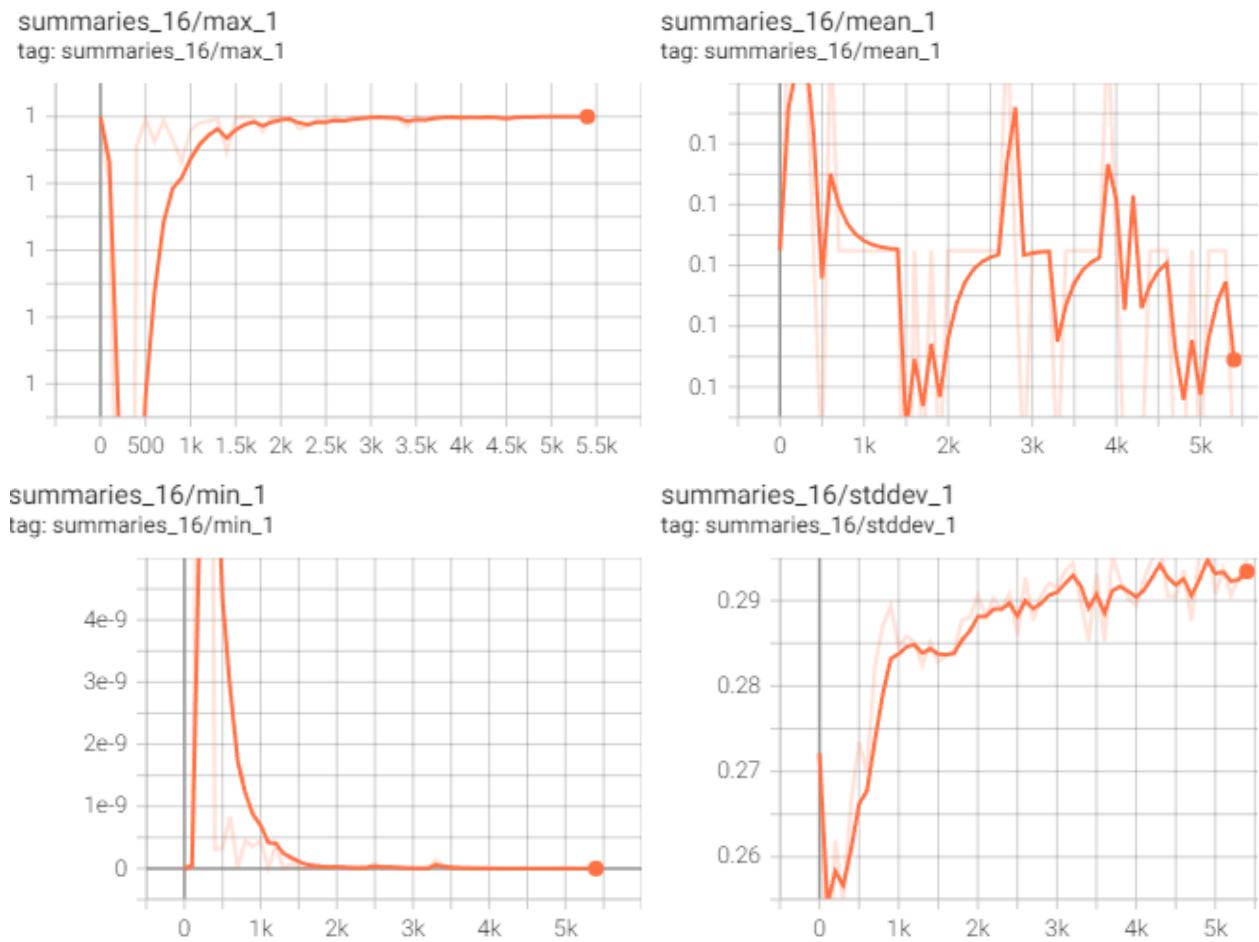


Figure 2\_36: y\_conv

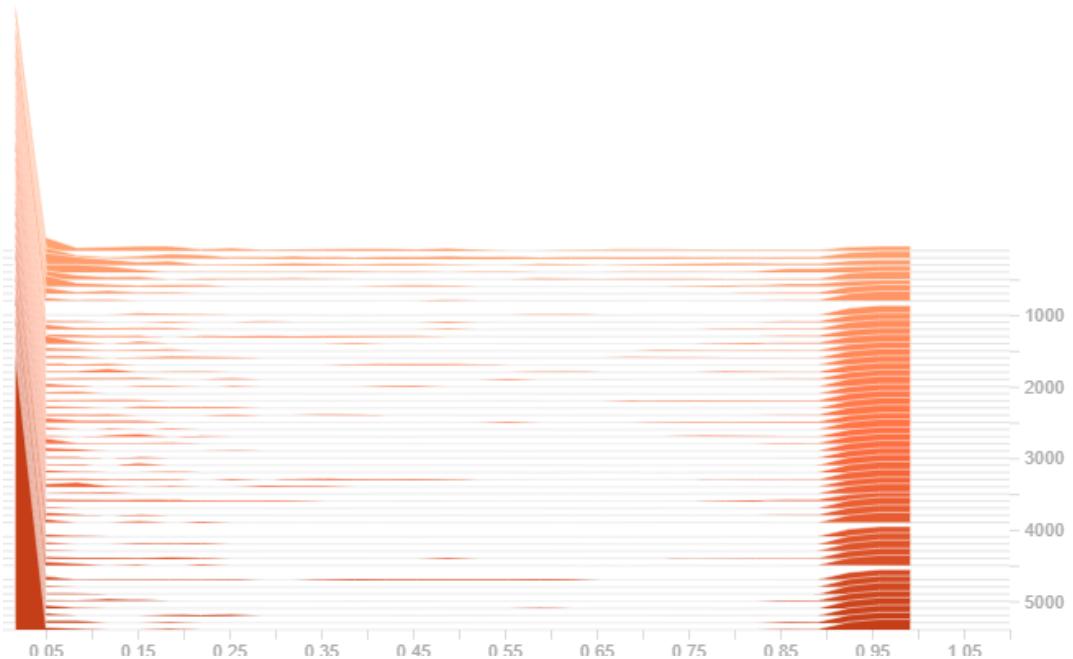


Figure 2\_37: y\_conv\_histogram

## 2.c

I change the relu to tanh for h\_conv1, h\_conv2, h\_fc1.

**Here** is the final output results:

step 0, training accuracy 0.06  
step 100, training accuracy 0.86  
step 200, training accuracy 0.84  
step 300, training accuracy 0.86  
step 400, training accuracy 0.88  
step 500, training accuracy 0.92  
step 600, training accuracy 0.96  
step 700, training accuracy 0.96  
step 800, training accuracy 0.98  
step 900, training accuracy 0.92  
step 1000, training accuracy 0.92  
step 1100, training accuracy 0.96  
step 1200, training accuracy 0.98  
step 1300, training accuracy 0.92  
step 1400, training accuracy 0.94  
step 1500, training accuracy 0.92  
step 1600, training accuracy 0.94  
step 1700, training accuracy 0.96  
step 1800, training accuracy 0.96  
step 1900, training accuracy 1  
step 2000, training accuracy 0.98  
step 2100, training accuracy 0.96  
step 2200, training accuracy 0.96  
step 2300, training accuracy 0.96  
step 2400, training accuracy 1  
step 2500, training accuracy 0.98  
step 2600, training accuracy 0.98  
step 2700, training accuracy 0.96  
step 2800, training accuracy 0.94  
step 2900, training accuracy 0.98  
step 3000, training accuracy 0.94  
step 3100, training accuracy 1  
step 3200, training accuracy 1  
step 3300, training accuracy 1

step 3400, training accuracy 1  
step 3500, training accuracy 1  
step 3600, training accuracy 1  
step 3700, training accuracy 1  
step 3800, training accuracy 0.96  
step 3900, training accuracy 1  
step 4000, training accuracy 0.98  
step 4100, training accuracy 0.98  
step 4200, training accuracy 1  
step 4300, training accuracy 0.98  
step 4400, training accuracy 0.98  
step 4500, training accuracy 0.98  
step 4600, training accuracy 0.98  
step 4700, training accuracy 0.98  
step 4800, training accuracy 0.98  
step 4900, training accuracy 1  
step 5000, training accuracy 1  
step 5100, training accuracy 1  
step 5200, training accuracy 0.98  
step 5300, training accuracy 1  
step 5400, training accuracy 1

test accuracy 0.9835

The training takes 879.099484 second to finish

The figure of scalars:

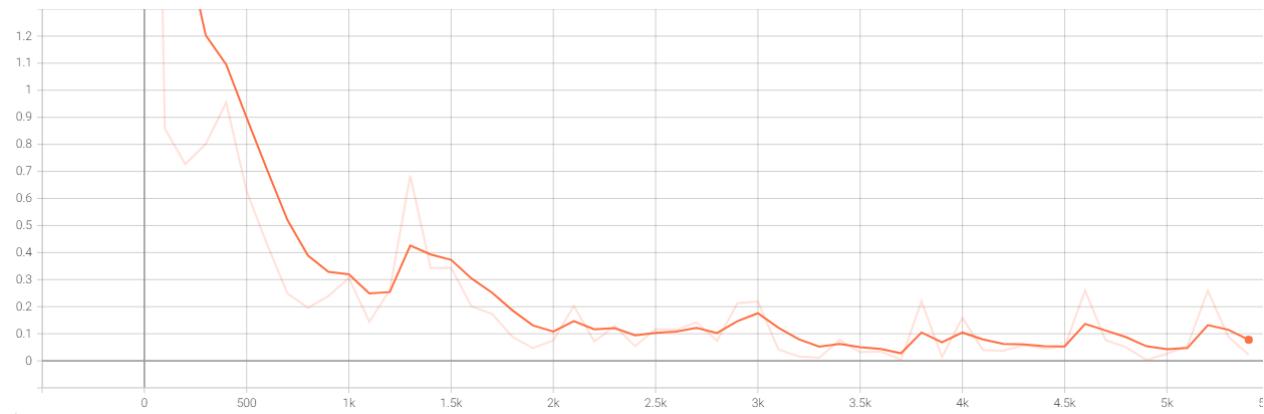
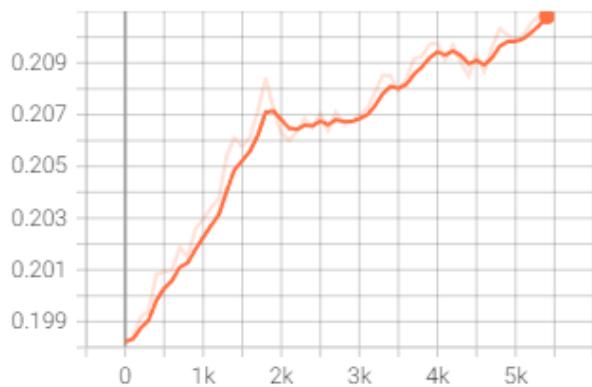


Figure 2\_38: scalars

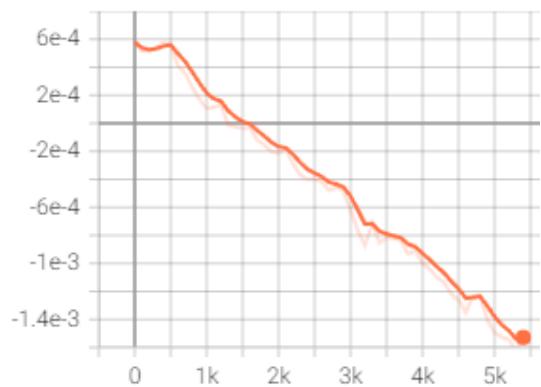
To use the function variable\_summaries, we run it on W\_conv1, b\_conv1, h\_conv1, h\_pool1, W\_conv2, b\_conv2, h\_conv2, h\_pool2, W\_fc1, b\_fc1, h\_pool2\_flat, h\_fc1, keep\_prob, h\_fc1\_drop, W\_fc2, b\_fc2, and y\_conv.

Here are the results:

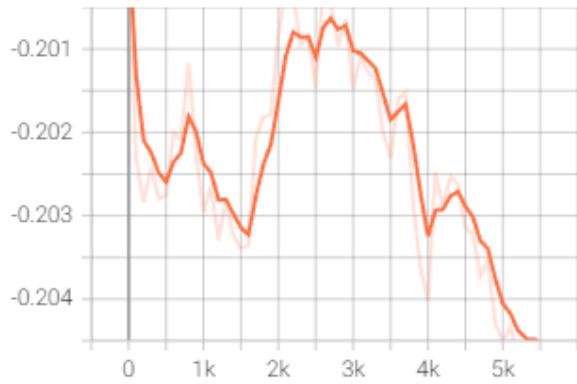
summaries/max\_1  
tag: summaries/max\_1



summaries/mean\_1  
tag: summaries/mean\_1



summaries/min\_1  
tag: summaries/min\_1



summaries/stddev\_1  
tag: summaries/stddev\_1

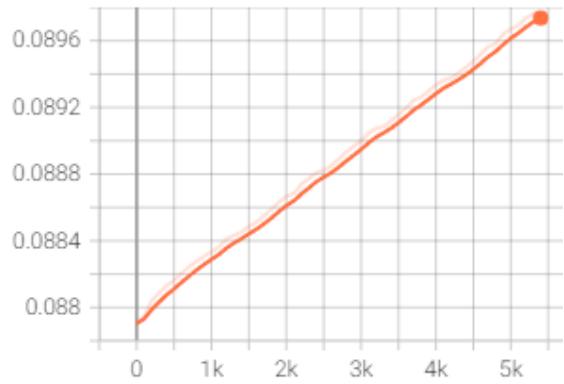


Figure 2\_39: W\_conv1

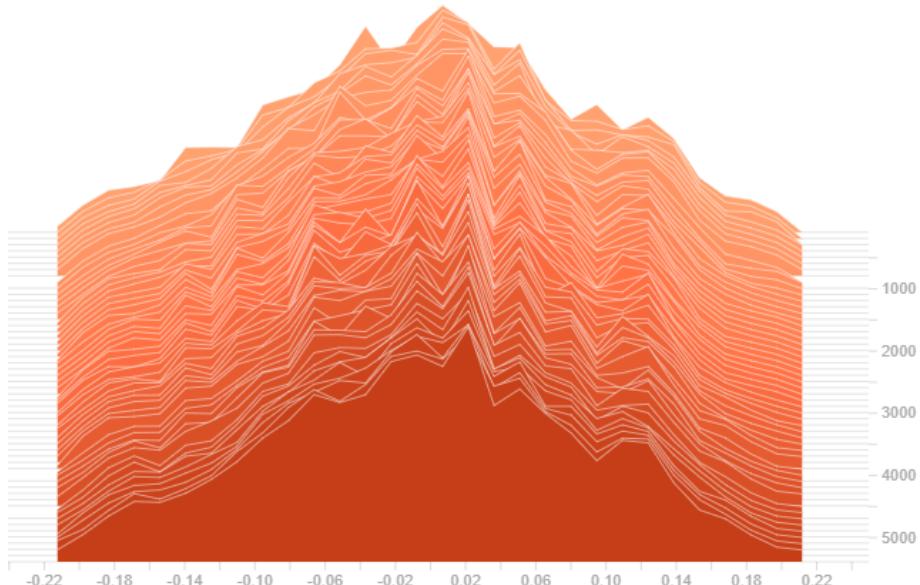


Figure 2\_40: W\_conv1\_histogram

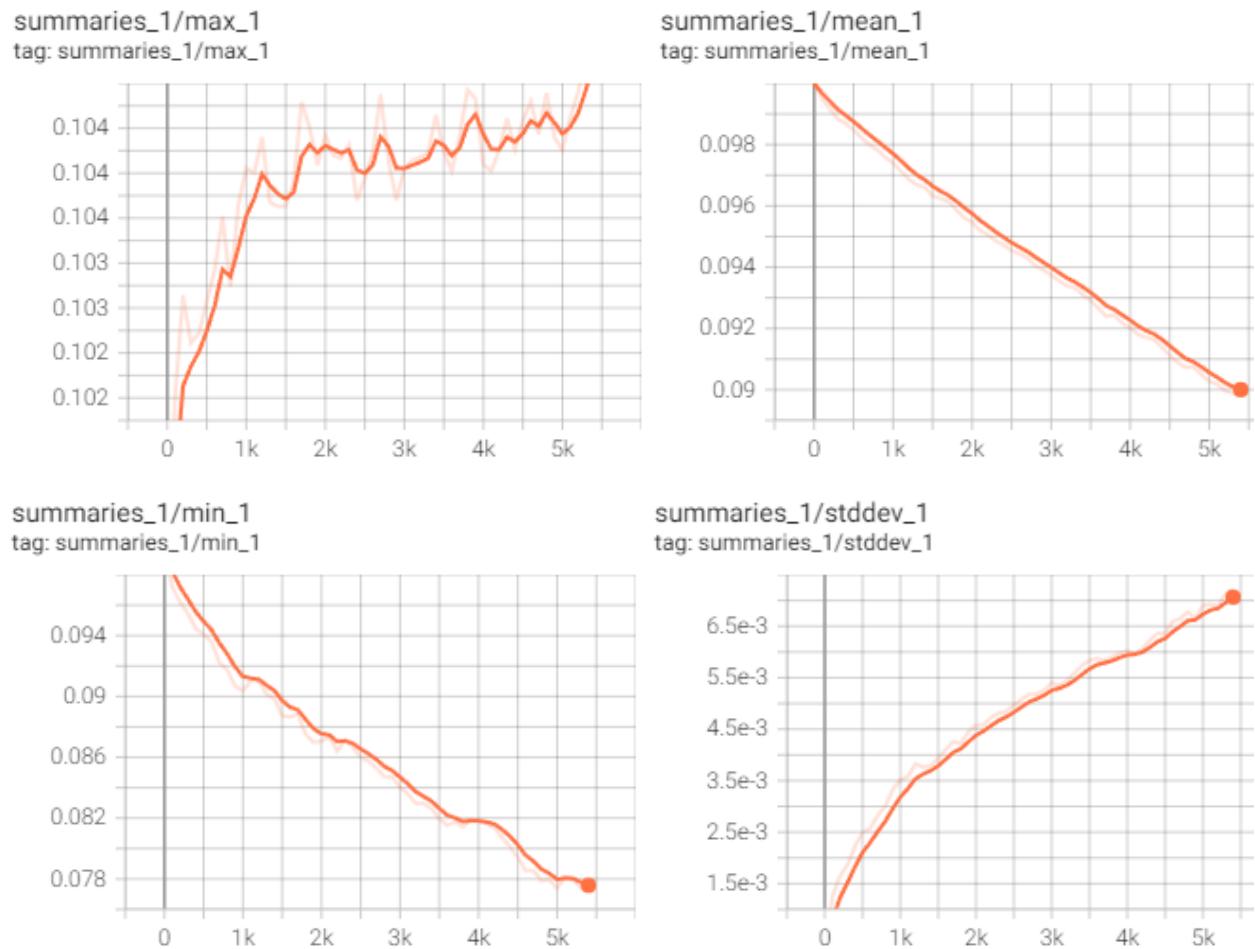


Figure 2\_41: b\_conv1

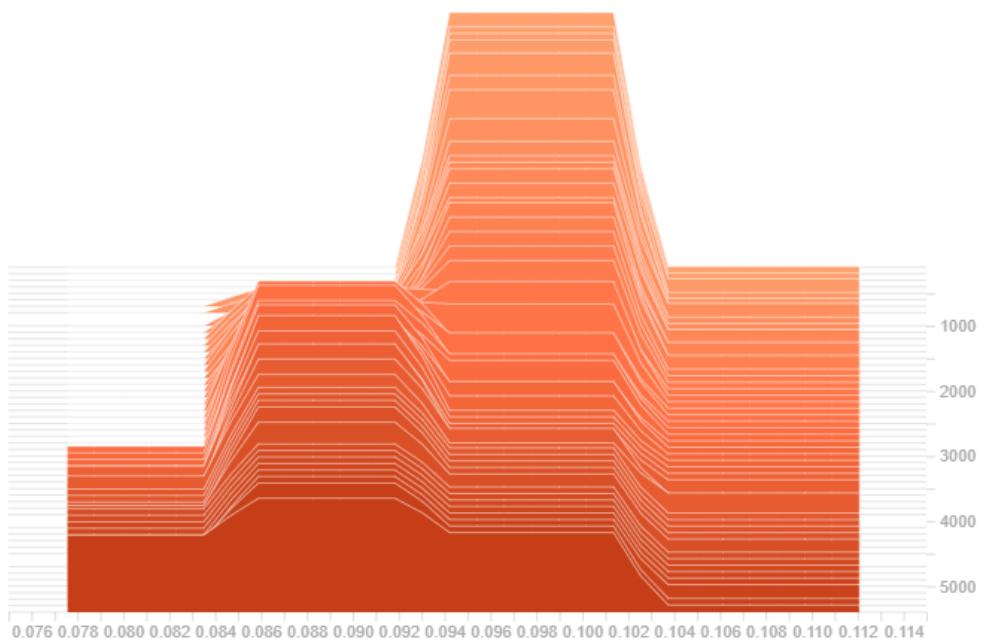


Figure 2\_42: b\_conv1\_histogram

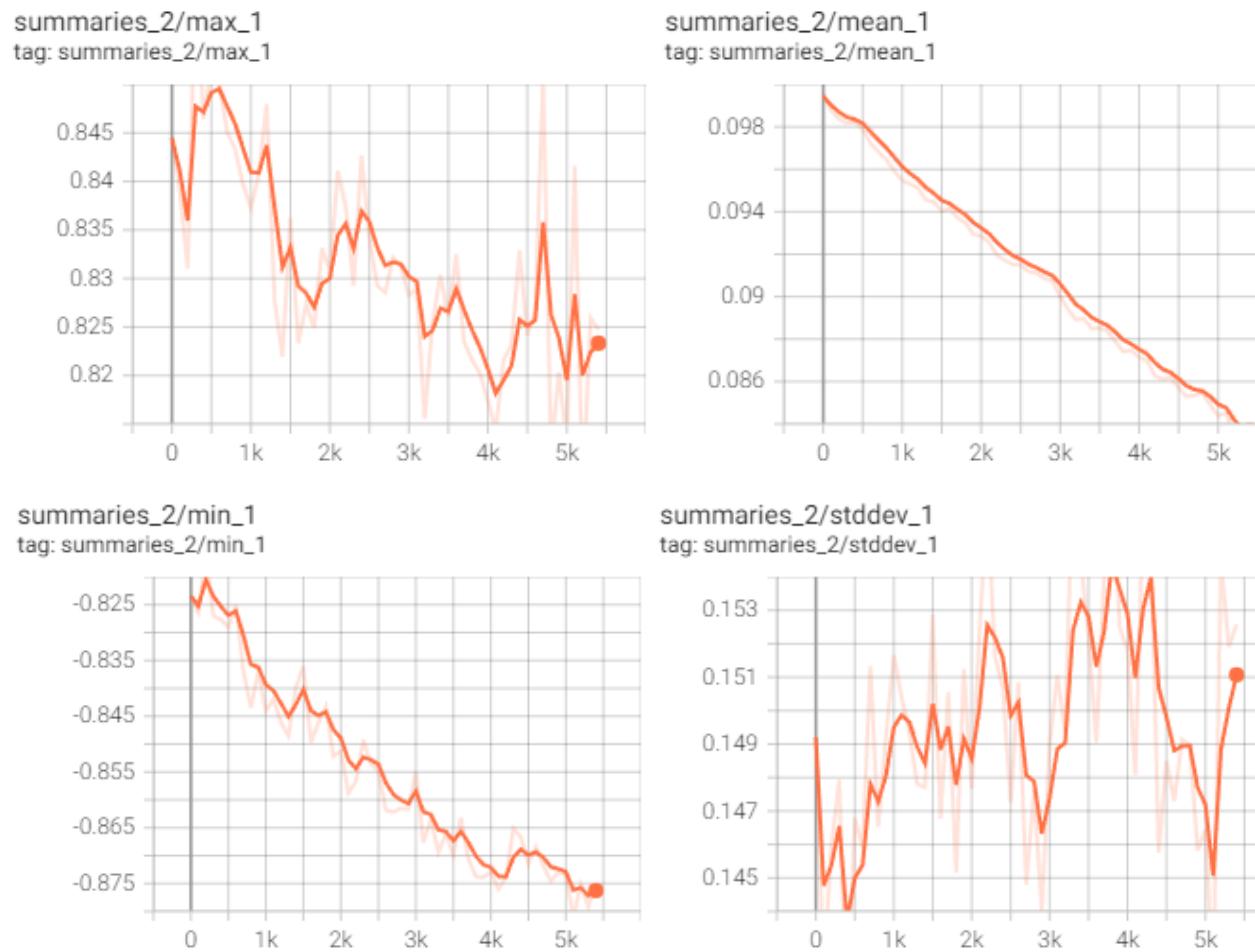


Figure 2\_43: h\_conv1

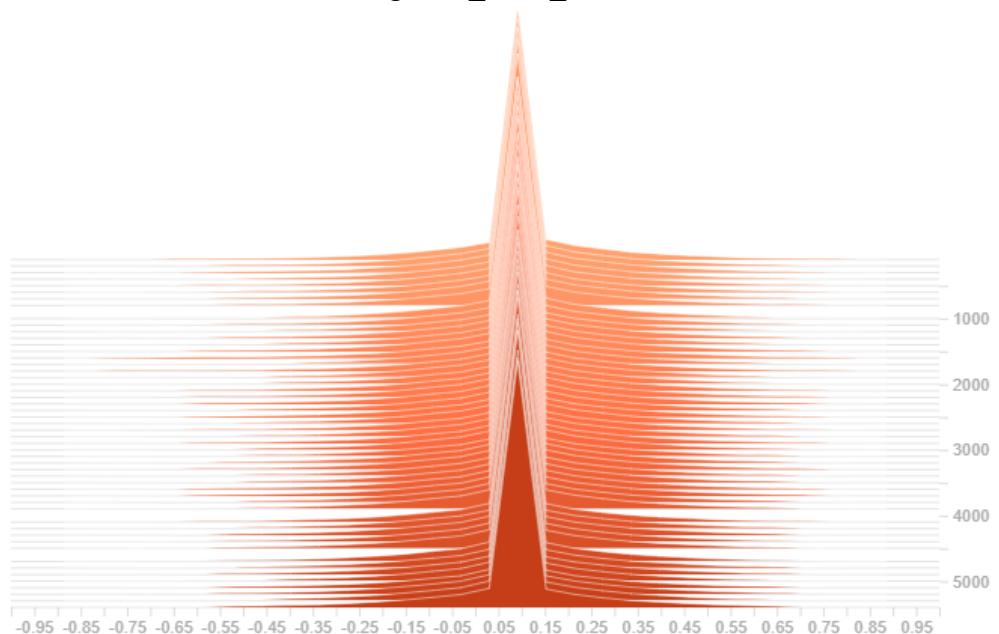
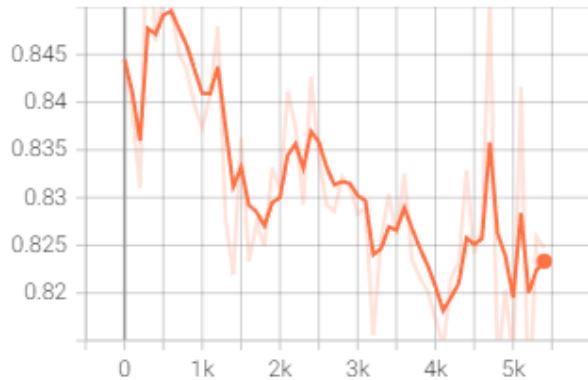
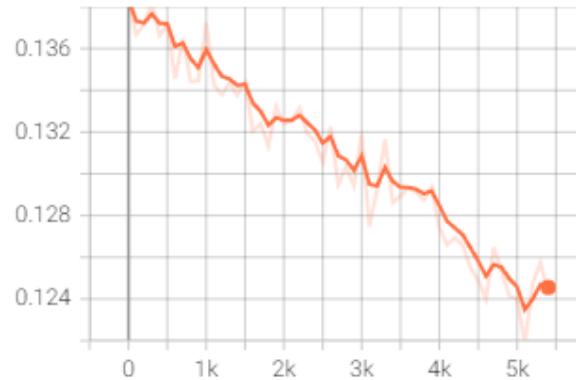


Figure 2\_44: h\_conv1\_histogram

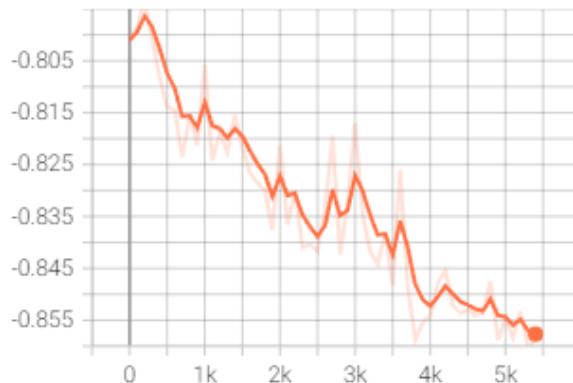
summaries\_3/max\_1  
tag: summaries\_3/max\_1



summaries\_3/mean\_1  
tag: summaries\_3/mean\_1



summaries\_3/min\_1  
tag: summaries\_3/min\_1



summaries\_3/stddev\_1  
tag: summaries\_3/stddev\_1

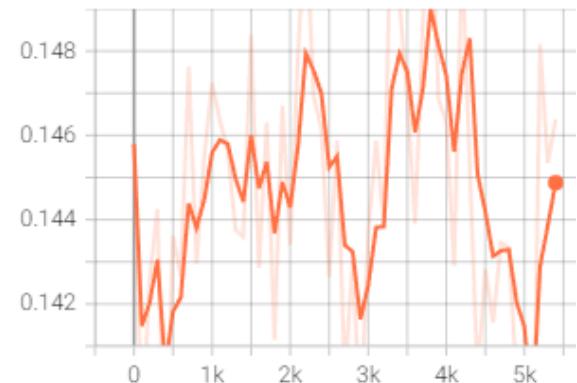


Figure 2\_45: h\_pool1

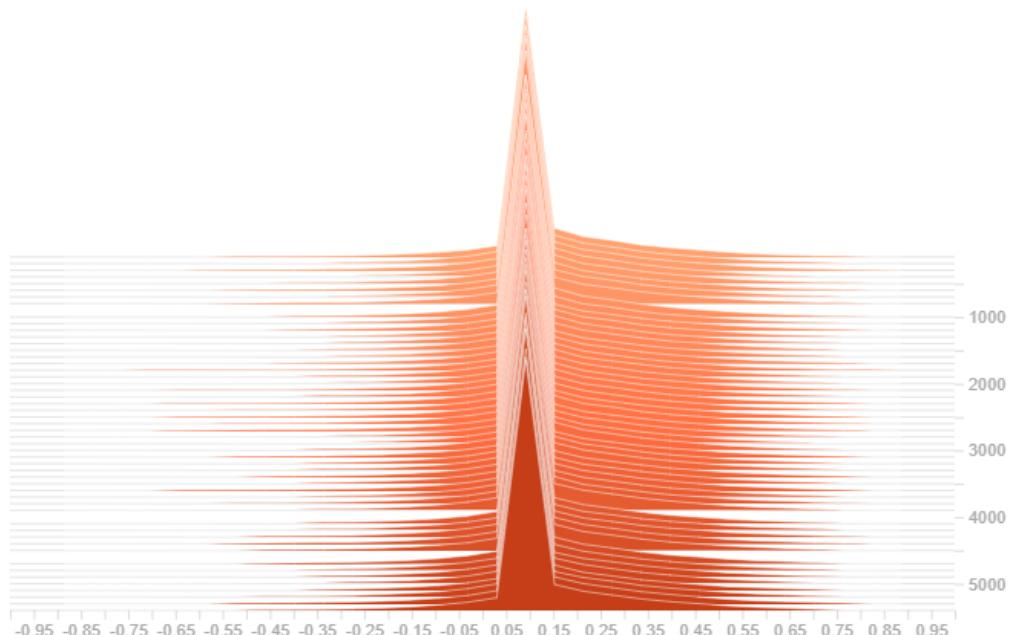
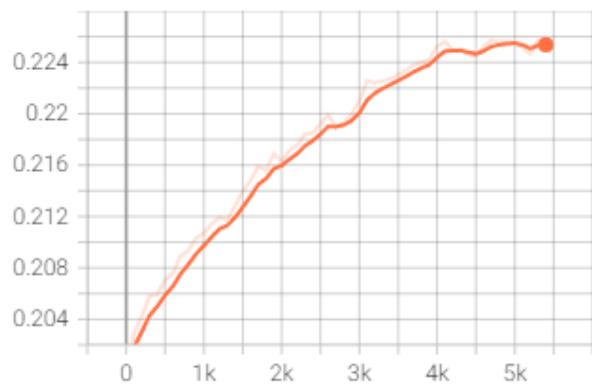
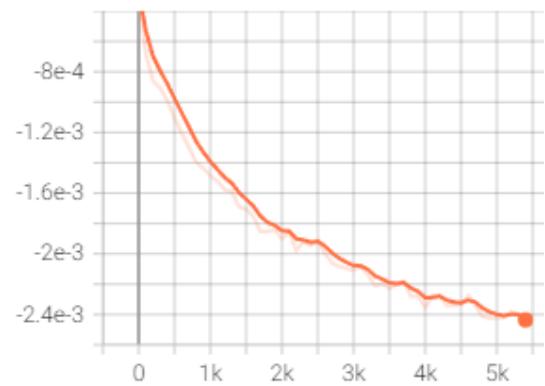


Figure 2\_46: h\_pool1\_histogram

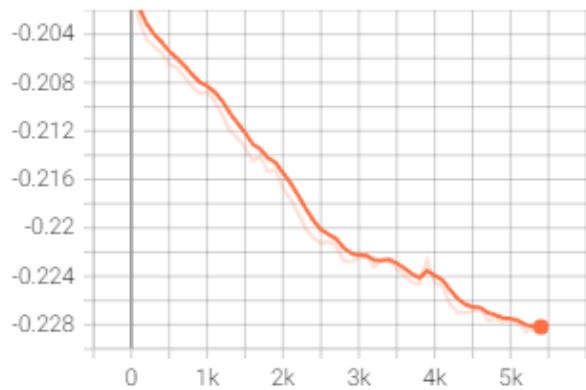
summaries\_4/max\_1  
tag: summaries\_4/max\_1



summaries\_4/mean\_1  
tag: summaries\_4/mean\_1



summaries\_4/min\_1  
tag: summaries\_4/min\_1



summaries\_4/stddev\_1  
tag: summaries\_4/stddev\_1

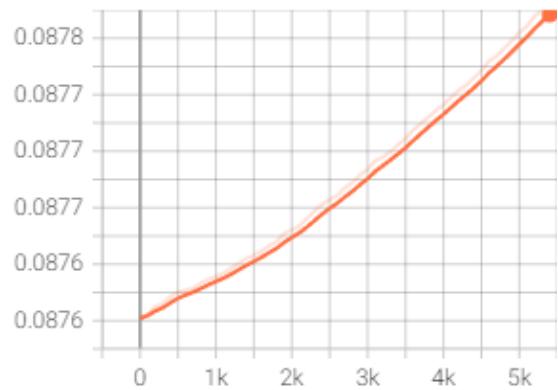


Figure 2\_47: W\_conv2

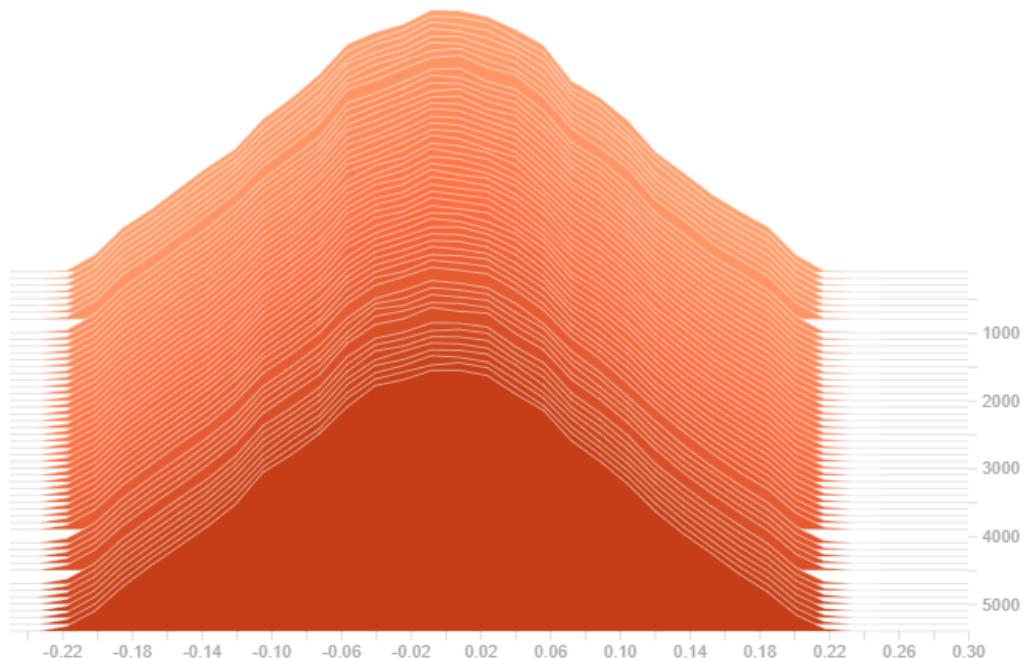


Figure 2\_48: W\_conv2\_histogram

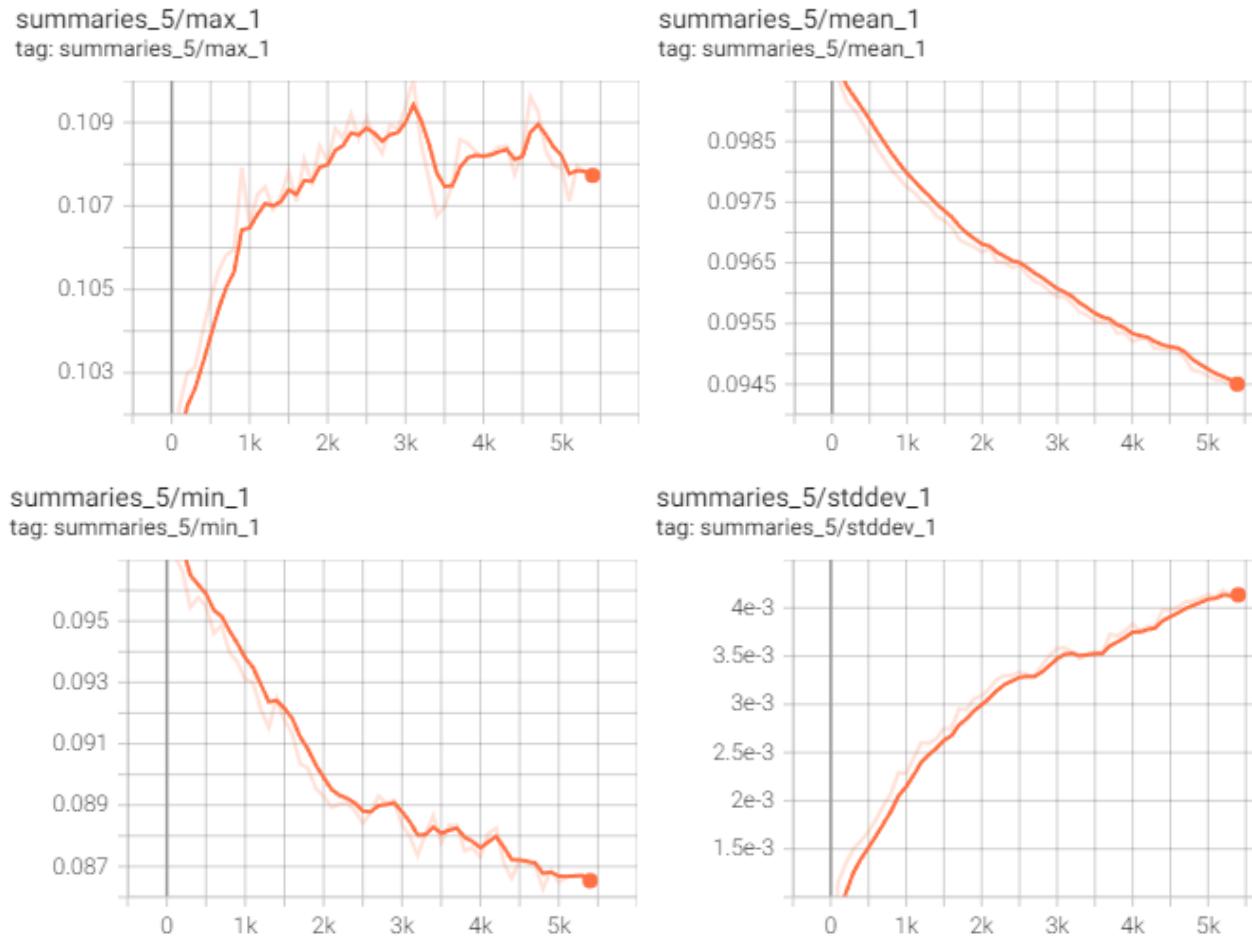


Figure 2\_49: b\_conv2

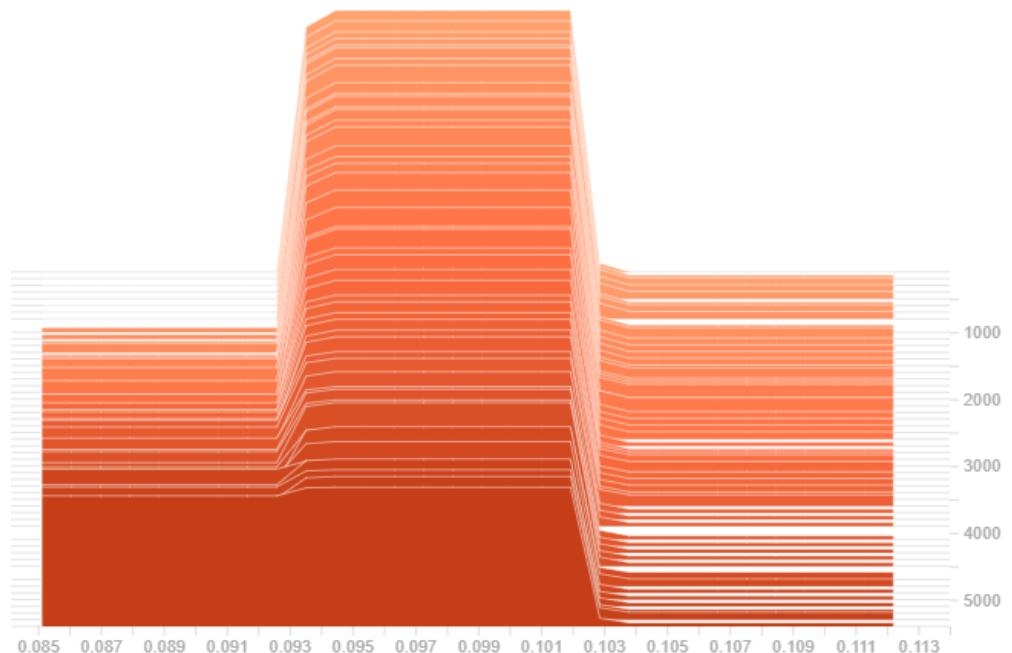


Figure 2\_50: b\_conv2\_histogram